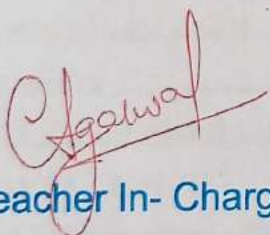


Thadomal Shahani Engineering College

Bandra (W.), Mumbai - 400 050.

❧ CERTIFICATE ❧

Certify that Mr./Miss Soumil Vivek Salvi
of IT Department, Semester V with
Roll No. 104 has completed a course of the necessary
experiments in the subject Advanced Dev Ops under my
supervision in the **Thadomal Shahani Engineering College**
Laboratory in the year 2023 - 2024


Teacher In- Charge

Head of the Department

Date 20/10/23

Principal

CONTENTS

SR. NO.	EXPERIMENTS	PAGE NO.	DATE	TEACHERS SIGN.
1)	To study & perform the setup of AWS EC2 service and launch an EC2 instance.		18/7/23	
2)	To study & perform the setup of AWS cloud9 service & launch a python program in cloud9.		25/7/23	
3)	To study AWS s3 service & create a bucket for hosting static web application.		1/8/23	
4)	To study codepipeline & deploy a web application using codepipeline.		2/8/23	
5)	To understand the kubernetes cluster architecture.		12/9/23	
6)	To understand the terraform life cycle & to build change and destroy AWS infrastructure using Terraform.		22/8/23	
7)	To perform static analysis on python programs using sonar Qube SAST process		29/8/23	
8)	To understand continuous monitoring using Nagios.		25/8/23	
9)	To understand AWS lambda function & create a lambda function using python to log "An image has been added"		31/8/23	
10)	To create AWS lambda function using python for adding data to dynamo DB d		5/9/23	
11)	Written Assignment - 1		23/9/23	
12)	Written Assignment - 2		3/10/23	

Approval
20/10/23

Approval
20/10/23

ASSIGNMENT-1

AIM : Study and create AWS EC2 instance

LO MAPPED : LO1

THEORY :

Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides secure, resizable compute capacity in the cloud. It is designed to make web-scale cloud computing easier for developers.

Amazon Elastic Compute Cloud (Amazon EC2) provides on-demand, scalable computing capacity in the Amazon Web Services (AWS) Cloud. Using Amazon EC2 reduces hardware costs so you can develop and deploy applications faster. You can use Amazon EC2 to launch as many or as few virtual servers as you need, configure security and networking, and manage storage. You can add capacity (scale up) to handle compute-heavy tasks, such as monthly or yearly processes, or spikes in website traffic. When usage decreases, you can reduce capacity (scale down) again.

Features of Amazon EC2

Amazon EC2 provides the following high-level features:

- Instances :

Virtual servers.

- Amazon Machine Images (AMIs):

Preconfigured templates for your instances that package the components you need for your server (including the operating system and additional software).

- Instance types :

Various configurations of CPU, memory, storage, networking capacity, and graphics hardware for your instances.

- Key pairs :

Secure login information for your instances. AWS stores the public key and you store the private key in a secure place.

- Instance store volumes :

Storage volumes for temporary data that is deleted when you stop, hibernate, or terminate your instance.

- Amazon EBS volumes :

Persistent storage volumes for your data using Amazon Elastic Block Store (Amazon EBS).

- Regions, Availability Zones, Local Zones, AWS Outposts, and Wavelength Zones :

Multiple physical locations for your resources, such as instances and Amazon EBS volumes.

- Security groups :

A virtual firewall that allows you to specify the protocols, ports, and source IP ranges that can reach your instances, and the destination IP ranges to which your instances can connect.

- Elastic IP addresses :

Static IPv4 addresses for dynamic cloud computing.

- Tags :

Metadata that you can create and assign to your Amazon EC2 resources.

- Virtual private clouds (VPCs) :

Virtual networks you can create that are logically isolated from the rest of the AWS Cloud. You can optionally connect these virtual networks to your own network.

STEPS-

1. LOGIN TO AWS ACCOUNT,
THEN SEARCH EC2.

23/07/2023, 22:59

Dashboard | EC2 Management Console

Services

More

New EC2 Experience

Tell us what you think

EC2 Dashboard

EC2 Global View

Events

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Resources

EC2 Global view

Refresh

You are using the following Amazon EC2 resources in the US East (N. Virginia) Region:

Instances (running)0

Auto Scaling Groups0

Dedicated Hosts0

Elastic IPs0

Instances0

Key pairs0

Load balancers0

Placement groups0

Security groups1

Snapshots0

Volumes0

Easily size, configure, and deploy Microsoft SQL Server Always On availability groups on AWS using the AWS Launch Wizard for SQL Server. [Learn more](#)

Account attributes

Refresh

Supported platforms

VPC

Default VPC

vpc-031db18c99362d042

Settings

EBS encryption

Zones

EC2 Serial Console

Default credit specification

Console experiments

Explore AWS

Close

Enable Best Price-

CloudShell

Feedback

Language

Privacy

Terms

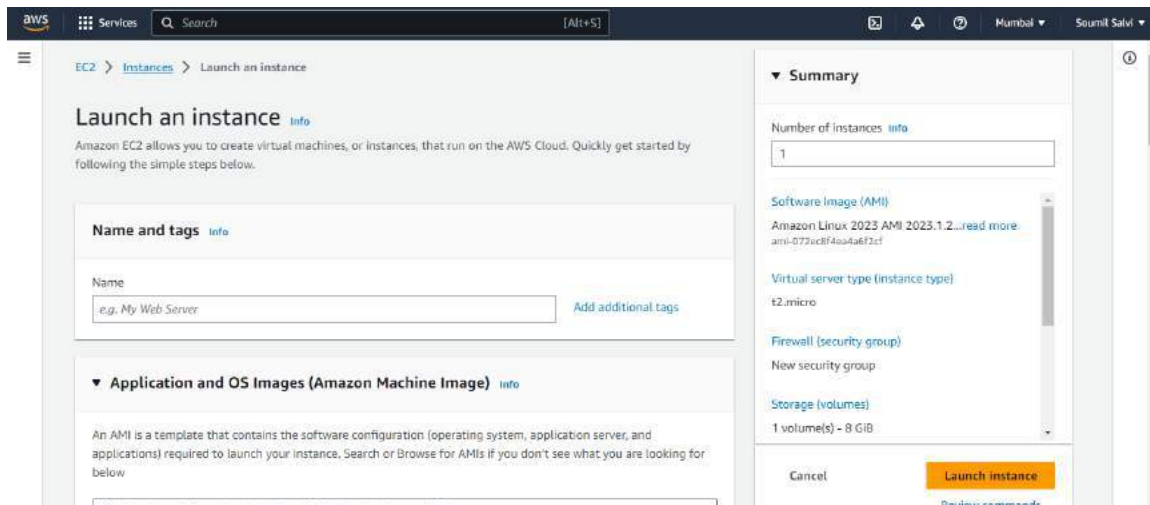
Cookie preferences

© 2023, Amazon Web Services India Private Limited or its affiliates.

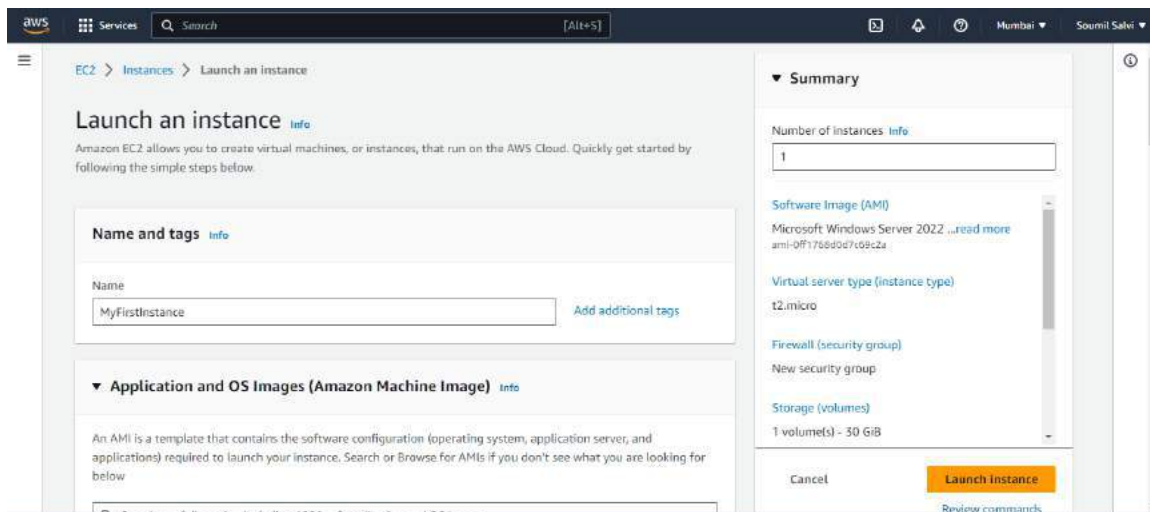
<https://us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#Home>

1/1

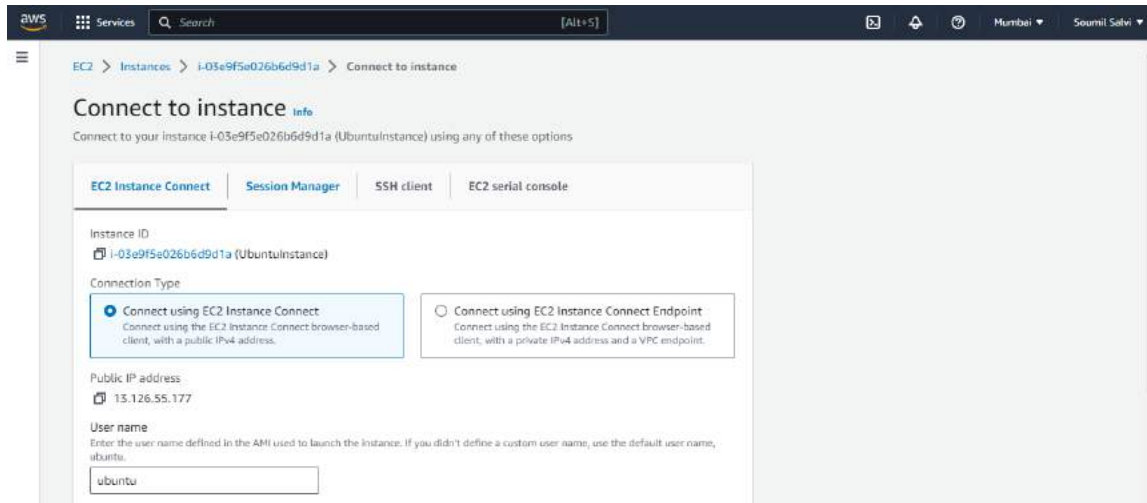
2. NOW CLICK ON LAUNCH / CREATE NEW INSTANCES.



3. Choose any machine you want to create here I am creating UBUNTU(free tier).

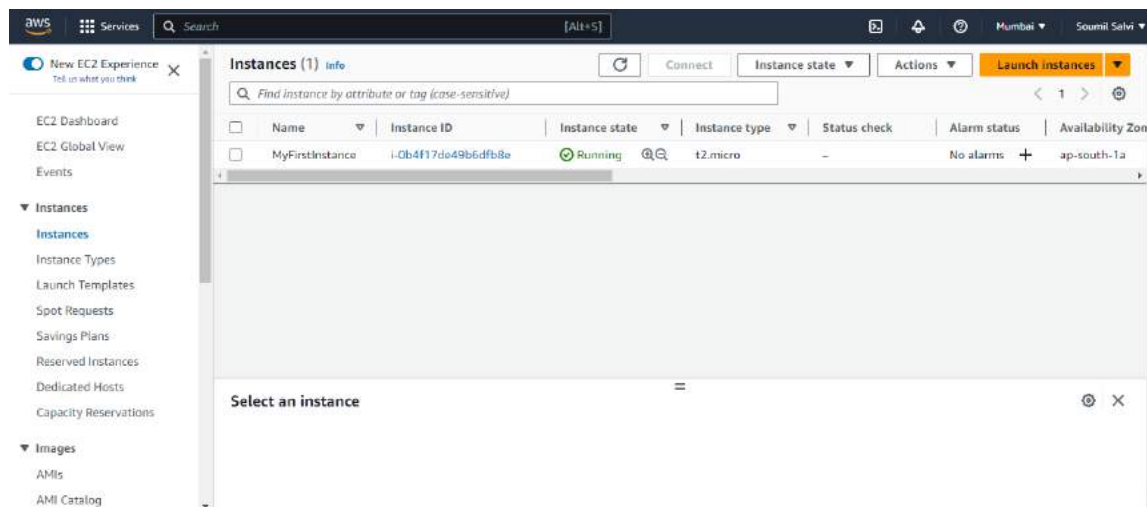


4. Now add security group ALL TRAFFIC ,
 PROTOCOL – ALL,
 SOURCE- ANYWHERE.



5. NOW WAIT TILL THE STATUS CHECK IS 2/2 and Instance is running.

Once Check is complete click on launch instances.



6. FOLLOW SOME BASIC LINUX COMMANDS AS SHOWN BELOW-

```
0 updates can be applied immediately.
Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo apt update

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-42-117:~$ ls
ubuntu@ip-172-31-42-117:~$
```

i-03e9f5e026b6d9d1a (UbuntuInstance)

PublicIPs: 13.126.55.177 PrivateIPs: 172.31.42.117

CONCLUSION : Hence learned and implemented the steps to create an EC2 machine.

ASSIGNMENT - 2

AIM : To understand the benefits of cloud Infrastructure and setup AWS Cloud 9 IDE, launch AWS Cloud9 IDE and perform collaboration demonstration.

LO MAPPED : LO1

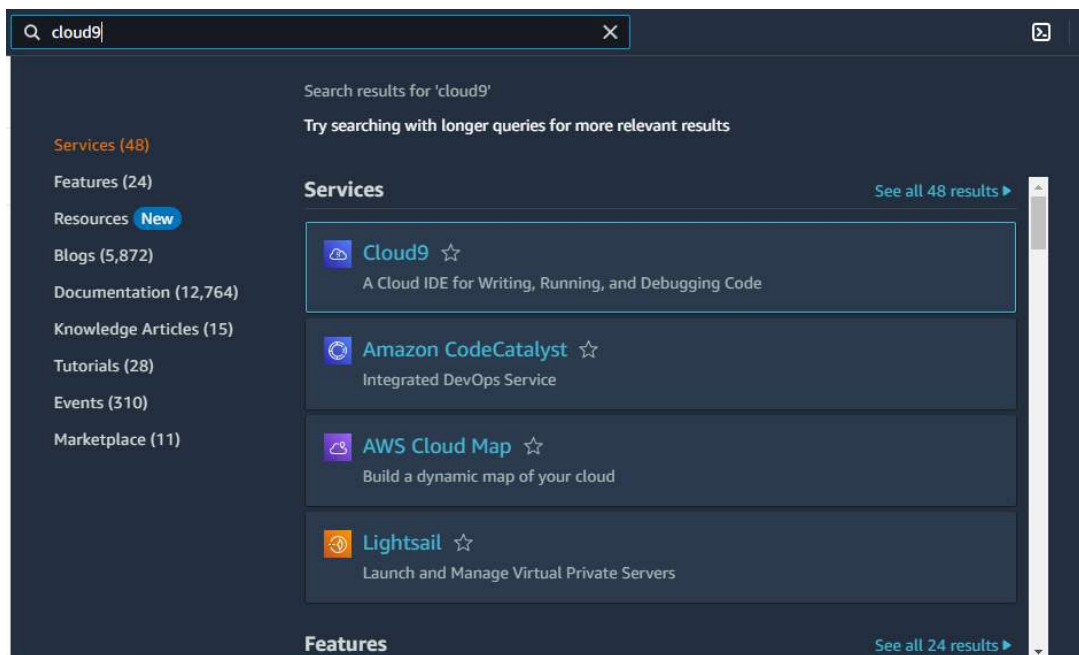
THEORY : AWS Cloud9 is an integrated development environment, or IDE.

The AWS Cloud9 IDE offers a rich code-editing experience with support for several programming languages and runtime debuggers, and a built-in terminal. It contains a collection of tools that you use to code, build, run, test, and debug software, and helps you release software to the cloud.

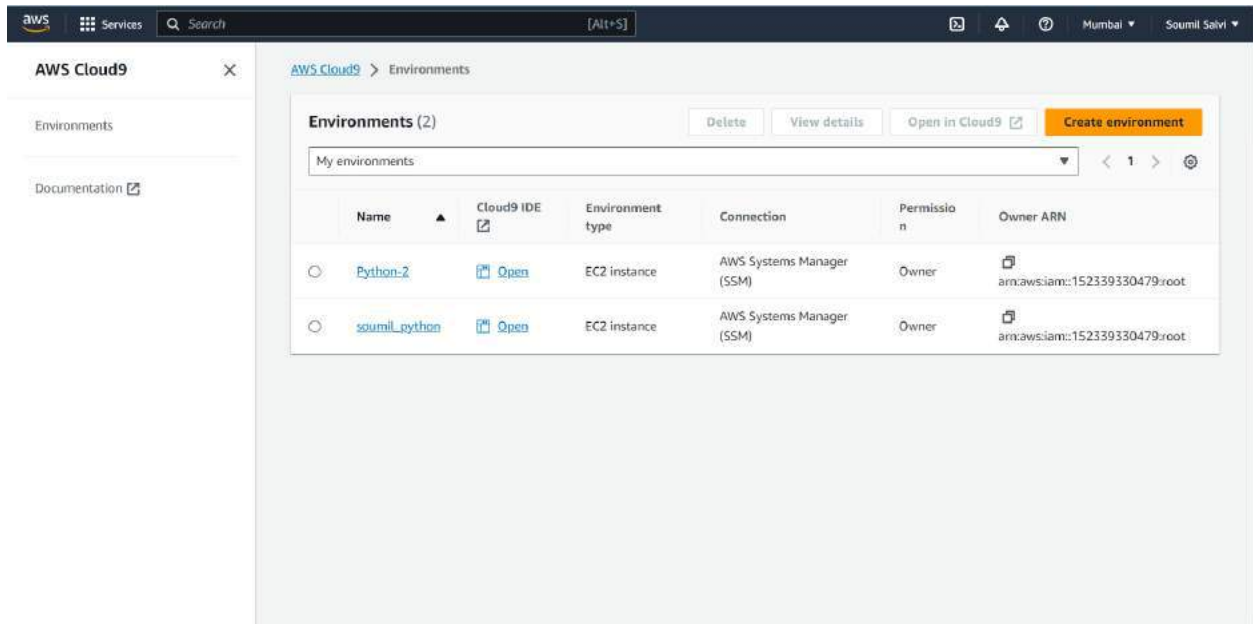
You access the AWS Cloud9 IDE through a web browser. You can configure the IDE to your preferences. You can switch color themes, bind shortcut keys, enable programming language-specific syntax coloring and code formatting, and more.

STEPS:

1) Search AWS Cloud in the AWS Dashboard.




2) Click on Cloud9, Enter your name and create your Cloud9 instance.



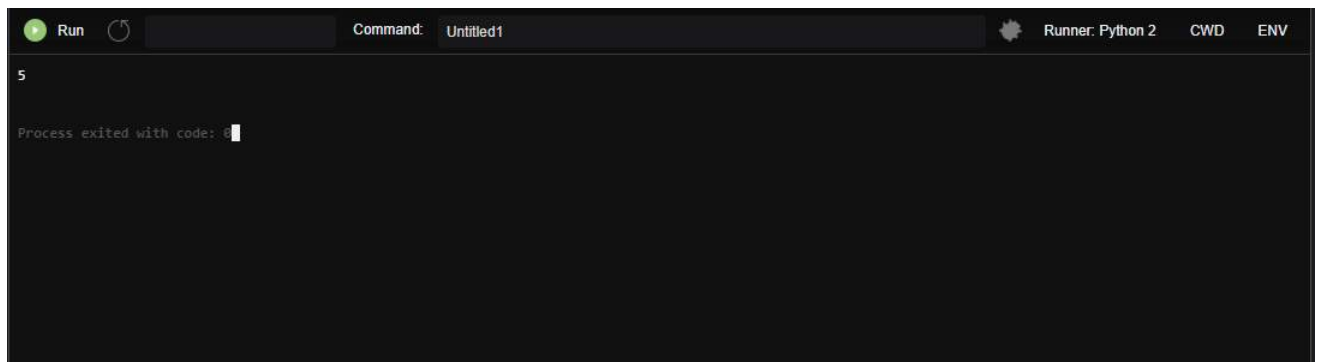
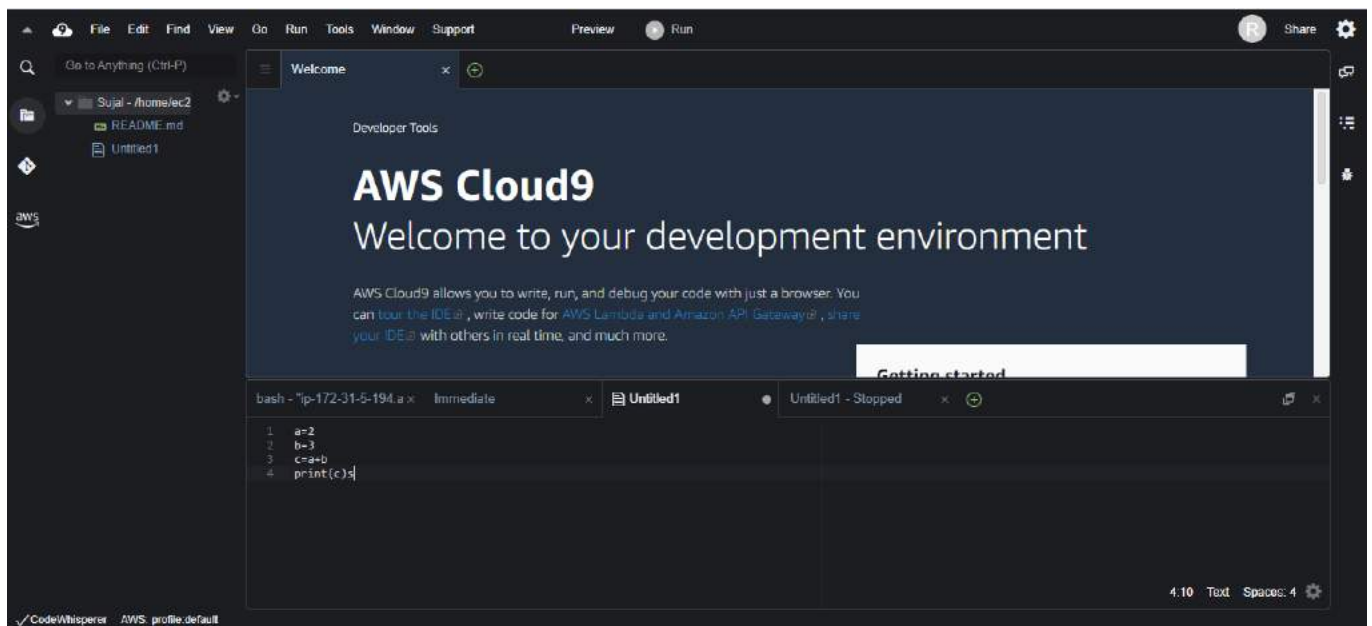
3) Click on open. This opens the Terminal. Start coding.

4) Finally, terminate the instance.

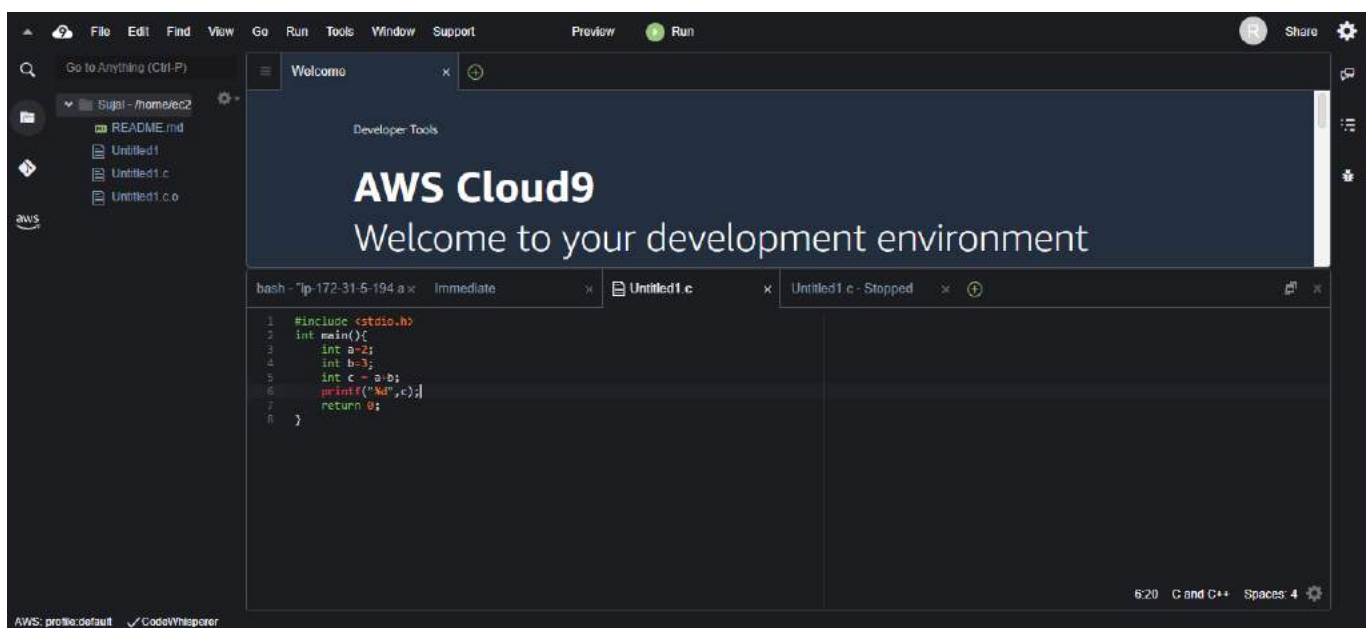


OUTPUT:

PYTHON:



C:



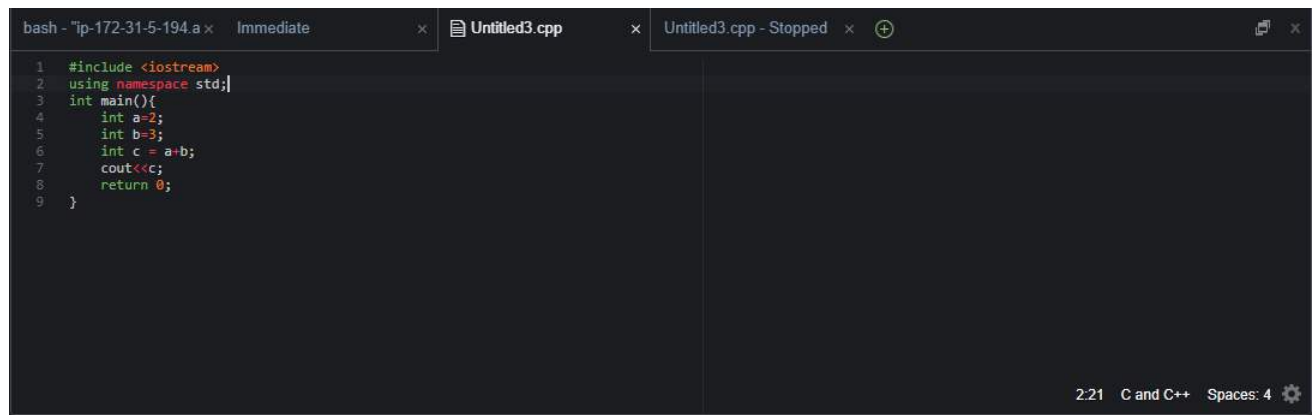
```
Run Command: Untitled1.c Runner: C CWD ENV
Running /home/ec2-user/environment/Untitled1.c
5
Process exited with code: 0
```

JAVA:

```
bash - "ip-172-31-5-194.a" x Immediate x exp.java x exp.java - Stopped x
1 import java.util.*;
2 public class exp{
3     public static void main(String []args){
4         int a=2;
5         int b=3;
6         System.out.println(a+b);
7     }
8 }
```

```
Run Command: exp.java Runner: Java CWD ENV
Building exp.java and running exp
5
Process exited with code: 0
```

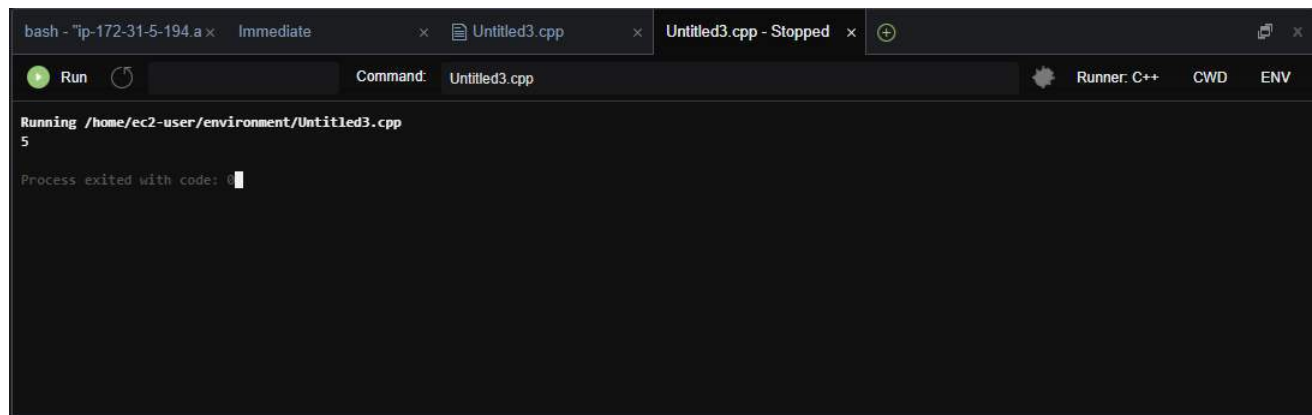
C++:



The screenshot shows the Cloud9 IDE editor with a dark theme. The top bar displays the terminal path 'bash - "ip-172-31-5-194.a x"', the file name 'Untitled3.cpp', and a status 'Stopped'. The editor contains a C++ program with the following code:

```
1  #include <iostream>
2  using namespace std;
3  int main(){
4      int a=2;
5      int b=3;
6      int c = a+b;
7      cout<<c;
8      return 0;
9  }
```

The bottom right corner of the editor shows the time '2:21', the language 'C and C++', and 'Spaces: 4'.



The screenshot shows the Cloud9 IDE terminal. The top bar displays the terminal path 'bash - "ip-172-31-5-194.a x"', the file name 'Untitled3.cpp', and a status 'Stopped'. The terminal has a 'Run' button and a 'Command' field containing 'Untitled3.cpp'. The runner is set to 'C++' with 'CWD' and 'ENV' options. The terminal output shows the program running and exiting with code 0.

```
Running /home/ec2-user/environment/Untitled3.cpp
5
Process exited with code: 0
```

CONCLUSION: AWS Cloud9 instance was created and Python, C, C++ and Java programs were performed using Cloud9 IDE.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 3

AIM : To study AWS S3 service and create a bucket for hosing static web application.

LO MAPPED : LO1,LO2

THEORY :

Amazon Simple Storage Service (Amazon S3) is a scalable object storage service provided by Amazon Web Services (AWS). It is designed to store and retrieve any amount of data from anywhere on the web. To understand the theory behind AWS S3 buckets, let's break down some key concepts:

- i. **Bucket:** A bucket is the fundamental container in S3. It's like a top-level folder or directory where you store your data. Each bucket has a globally unique name, meaning no two buckets can have the same name across all of AWS. For example, you might have a bucket named "mycompanydata."
- ii. **Objects:** Objects are the data files stored within S3 buckets. These can be anything from text files and images to videos and databases. Each object in a bucket has a unique key that distinguishes it from other objects in the same bucket. The combination of the bucket name and object key forms a unique identifier for the object within S3. For example, an object might have a key like "user-profiles/user123.jpg."
- iii. **Data Consistency:** S3 provides strong read-after-write consistency for all objects in your bucket. This means that once you write an object to an S3 bucket, you can immediately read it, and you will always get the latest version of the object.
- iv. **Data Durability:** S3 is designed to provide 99.999999999% (11 nines) durability, which means that your data is highly resilient and protected against data loss. This durability is achieved through data replication across multiple Availability Zones within a region.

- v. **Data Availability:** S3 is designed for high availability, ensuring that your data is accessible when you need it. Availability is typically in the range of 99.99%, and it can be even higher depending on how you configure your bucket.
- vi. **Storage Classes:** S3 offers various storage classes, each with different characteristics and pricing. These storage classes include Standard, Intelligent-Tiering, Glacier, and more. You can choose the appropriate storage class for your specific use case based on factors like data access frequency and retrieval time requirements.
- vii. **Access Control:** You can control access to your S3 buckets and objects using AWS Identity and Access Management (IAM) policies and Access Control Lists (ACLs). This allows you to specify who can upload, download, or delete objects within your buckets.
- viii. **Data Lifecycle Management:** S3 allows you to define data lifecycle policies to automatically transition objects between storage classes or delete them after a certain period. This can help you optimize storage costs.
- ix. **Versioning:** S3 supports versioning, which means you can keep multiple versions of an object in a bucket. This is useful for data recovery and maintaining a version history.
- x. **Security:** S3 offers features like encryption at rest and in transit, which help protect your data. You can also configure bucket policies to restrict access to specific IP addresses or VPCs.
- xi. **Event Notifications:** You can set up event notifications on S3 buckets to trigger actions (e.g., invoking a Lambda function) when specific events occur, such as object creation or deletion.
- xii. **Cross-Region Replication:** You can replicate your S3 data to other AWS regions for disaster recovery, compliance, or low-latency access to data in different geographical locations.

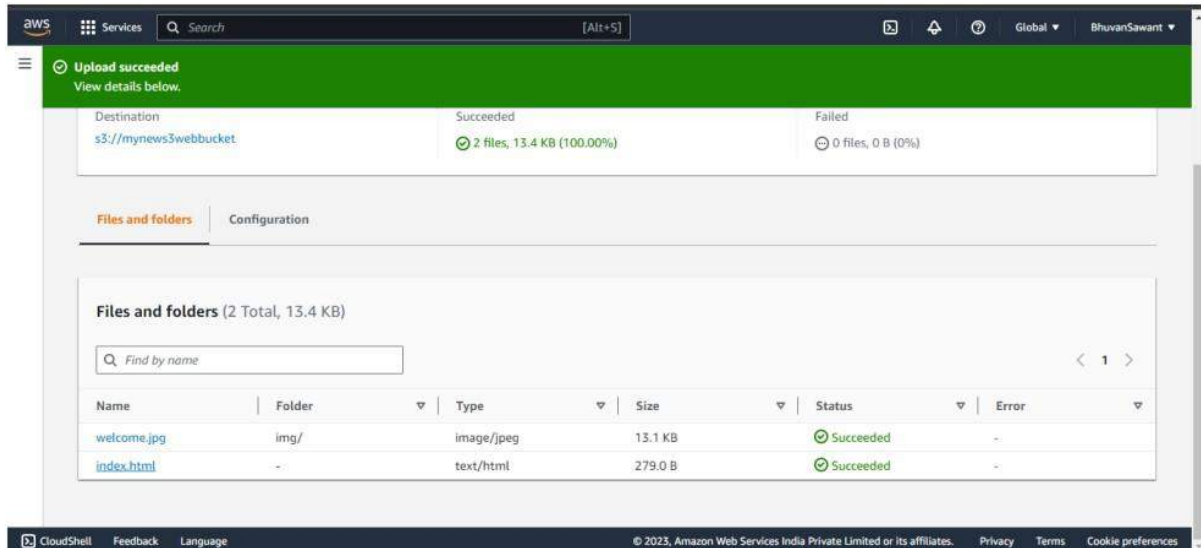
In summary, AWS S3 buckets are versatile, highly durable, and scalable containers for storing objects of various types. They provide robust data management capabilities, security features, and access control mechanisms, making them a fundamental component of many cloud-based applications and data storage solutions. Understanding these concepts is essential for effectively using S3 to store and manage your data in the AWS cloud.

Use Cases:

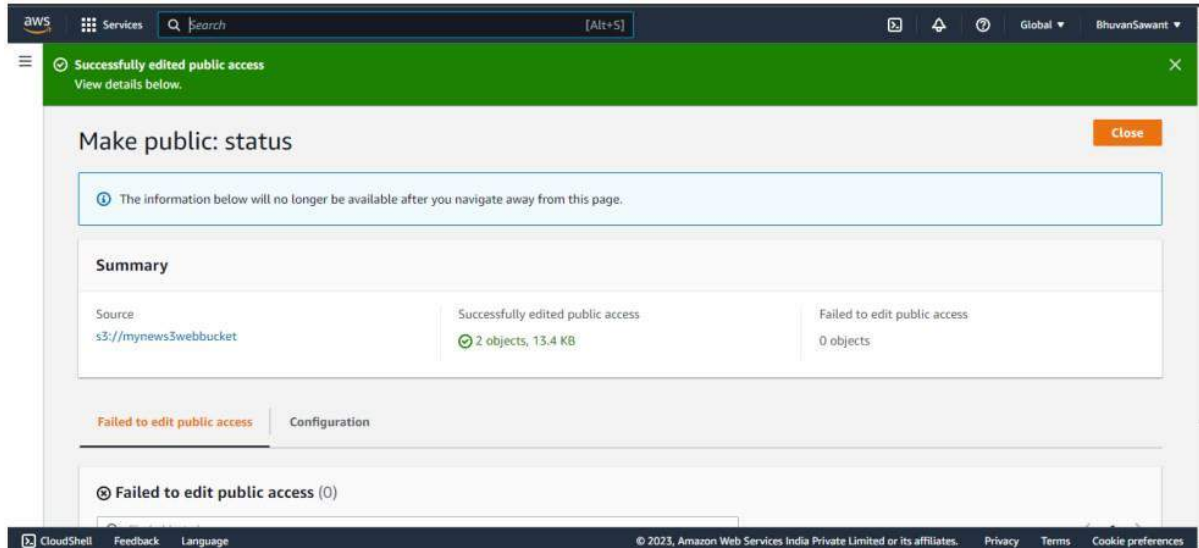
- i. **Data Storage:** S3 is primarily used for storing and archiving data. It can store virtually any type of data, including documents, images, videos, backups, logs, and databases.
- ii. **Static Website Hosting:** You can host static websites directly on S3 by configuring it to serve HTML, CSS, JavaScript, and other web assets. This is a cost-effective way to host websites with high availability and scalability.
- iii. **Data Backup and Recovery:** Many organizations use S3 as a backup and disaster recovery solution. It offers high durability and can be configured to automatically backup data from on-premises servers or other AWS resources.
- iv. **Data Lakes:** S3 is often used as a foundation for building data lakes, which are centralized repositories for storing and analyzing large volumes of structured and unstructured data. It can integrate with various analytics and data processing services in AWS.
- v. **Content Delivery:** Amazon S3 can be integrated with Amazon CloudFront (AWS's content delivery network) to deliver content like images, videos, and other media files to users around the world with low latency and high data transfer speeds.
- vi. **Big Data and Analytics:** S3 is commonly used as a data store for big data and analytics platforms like Amazon EMR, Amazon Redshift, and AWS Glue. It allows you to store vast amounts of data and access it for analysis.
- vii. **IoT Data Storage:** Internet of Things (IoT) devices can send data directly to S3 for storage and analysis, enabling IoT applications to collect and process large volumes of sensor data.
- viii. **Log Storage and Analysis:** Many organizations use S3 to store log files generated by applications and infrastructure. Services like Amazon Athena can be used to analyze these logs directly in S3.

Steps :-


- i. First login to your AWS account.



- ii. Then create a new S3 bucket.



- iii. After that host a static web page using the S3 bucket.



Website for AWS

CONCLUSION : In this assignment we learnt about AWS S3 bucket and hosted a static web page using the S3 bucket.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 4

AIM : To study AWS Code Pipeline and deploy web application using Code Pipeline.

LO MAPPED : LO1,LO2

THEORY :

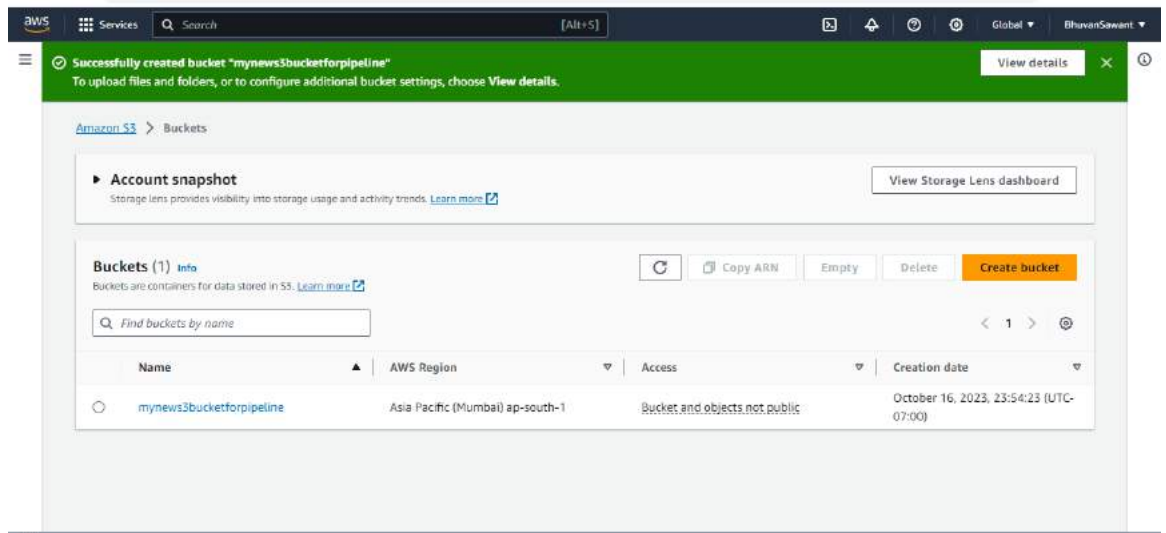
AWS CodePipeline is a continuous integration and continuous delivery (CI/CD) service provided by Amazon Web Services (AWS).

The key aspects of AWS CodePipeline:

Overview:

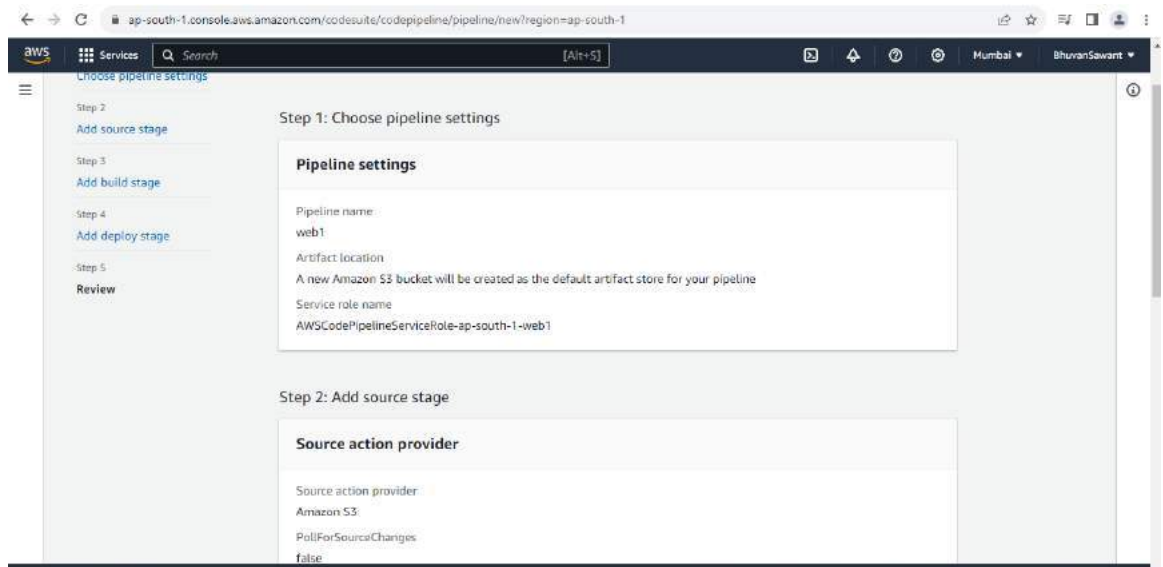
1. CI/CD Workflow:

- AWS CodePipeline facilitates the automation of the build, test, and deployment phases of the release process. It allows you to define a series of stages, each of which can represent a phase in your release pipeline.



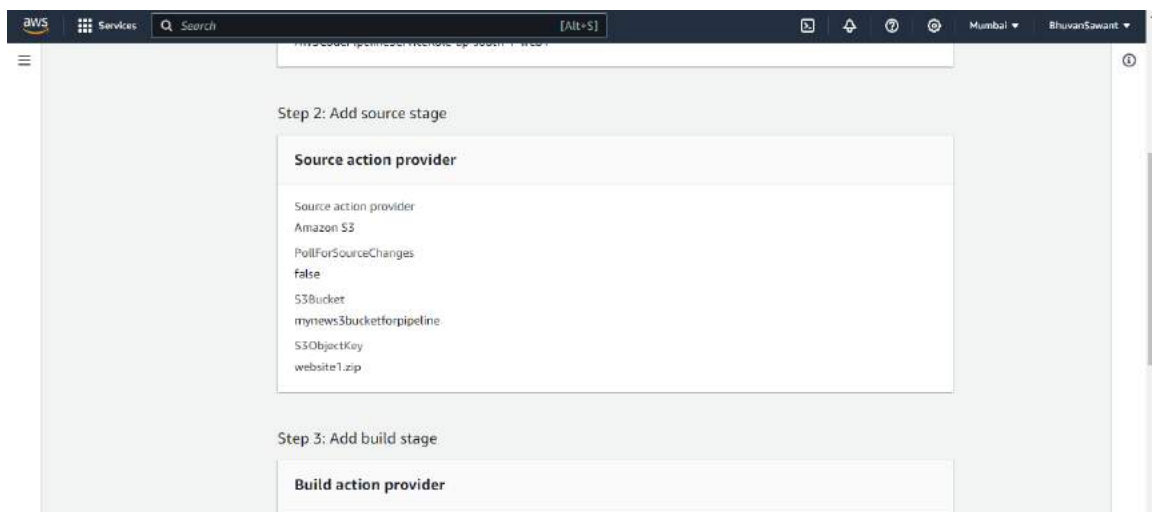
2. Integration with Other AWS Services:

- CodePipeline integrates with various AWS services, such as AWS CodeBuild for building applications, AWS CodeDeploy for automating deployments, and AWS Lambda for running custom actions.



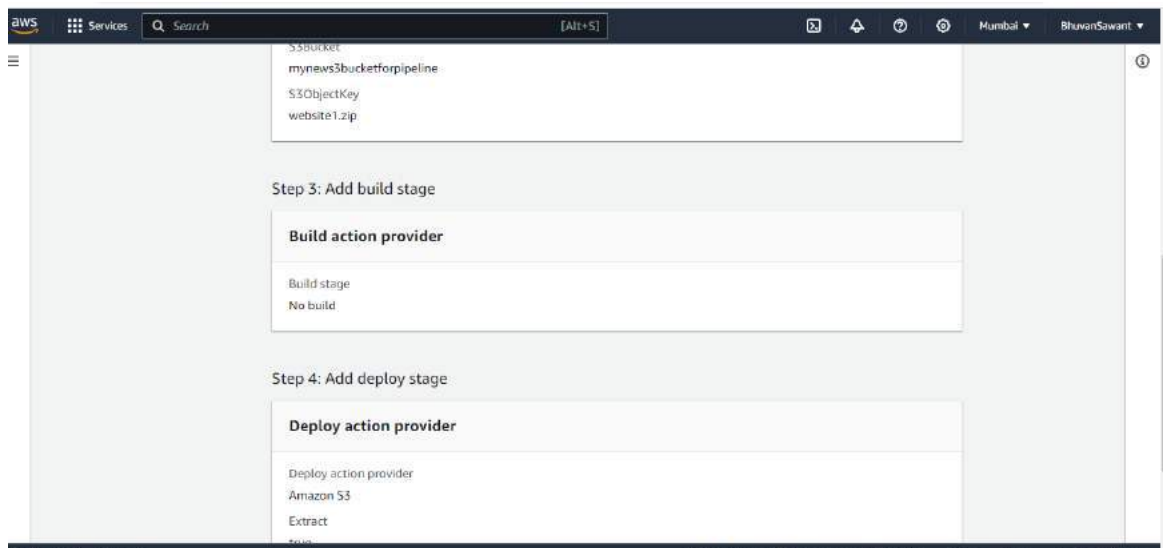
3. Pipeline Execution:

- Pipelines consist of a series of stages, and each stage can have one or more actions. Actions represent a task, such as source code retrieval or deployment to a specific environment.



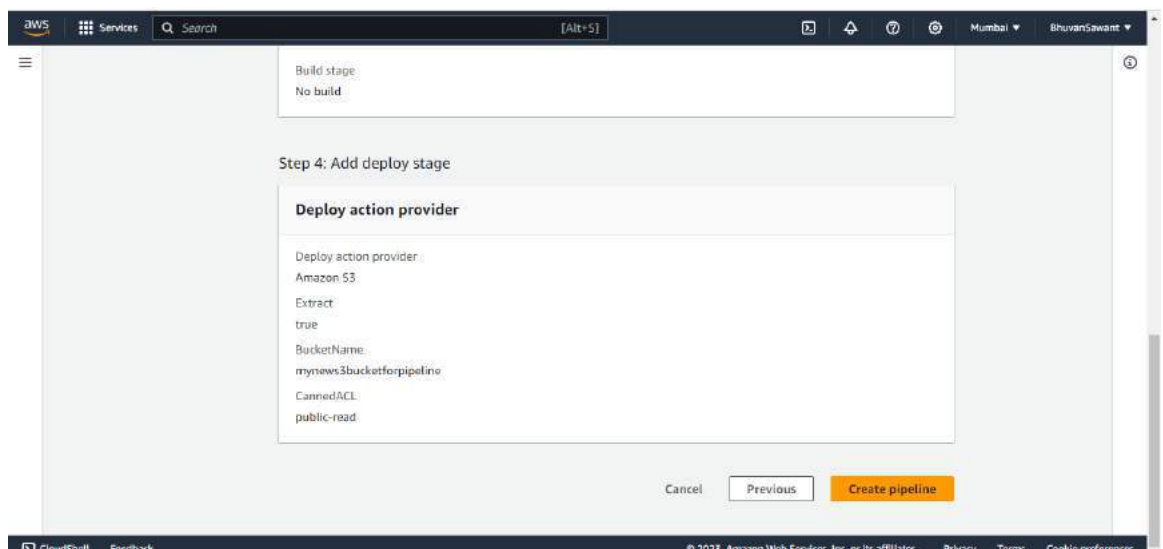
4. Source Providers:

- CodePipeline supports integration with various source code repositories, including AWS CodeCommit, GitHub, and Amazon S3.



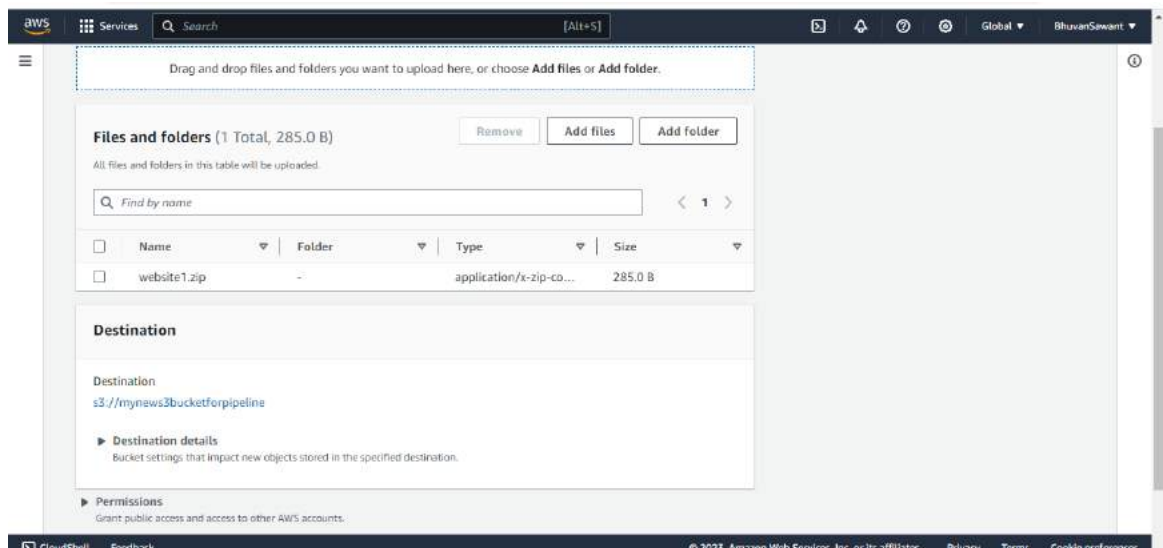
5. Artifact Management:

- CodePipeline uses artifacts to store the files and data needed for each action in a pipeline. Artifacts can be passed between stages to ensure consistency in the deployment process.



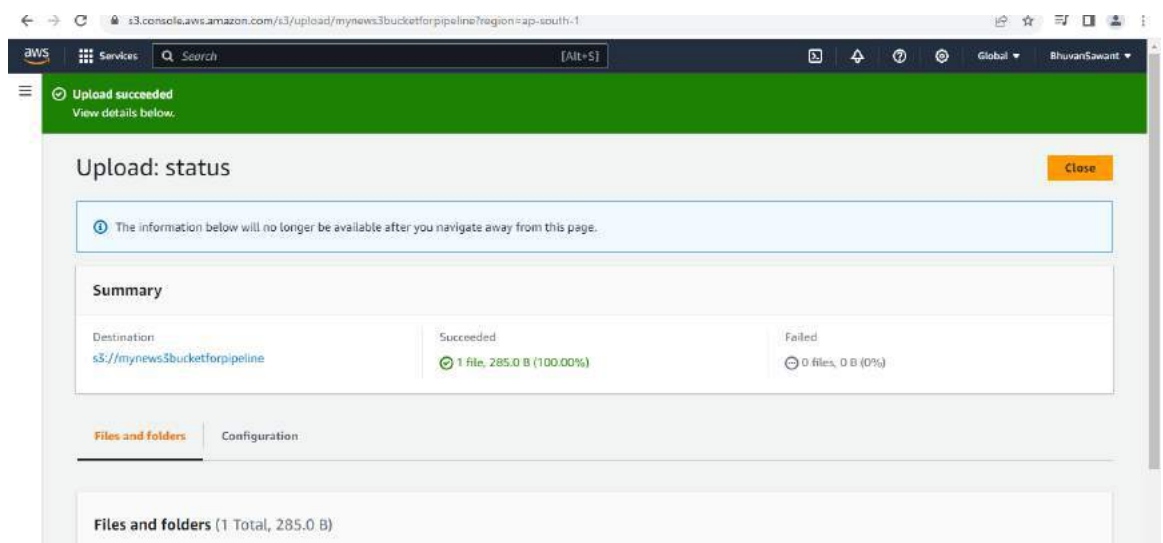
6. Integration with Third-Party Tools:

- Besides AWS services, CodePipeline supports integration with third-party tools. This is achieved through custom actions, which allow you to use external tools and scripts in your pipeline.



7. Pipeline Visualizations:

- CodePipeline provides a visual representation of your release process, making it easy to understand and monitor the status of each stage and action.



Key Concepts:

1. Pipeline:

- A pipeline is a series of stages that represents your release process. Each stage can contain one or more actions.

2. Stage:

- A stage is a logical unit in a pipeline, representing a phase in the release process. Stages are executed sequentially.

3. Action:

- An action represents a task within a stage. Actions can include tasks such as building code, deploying to a test environment, or running tests.

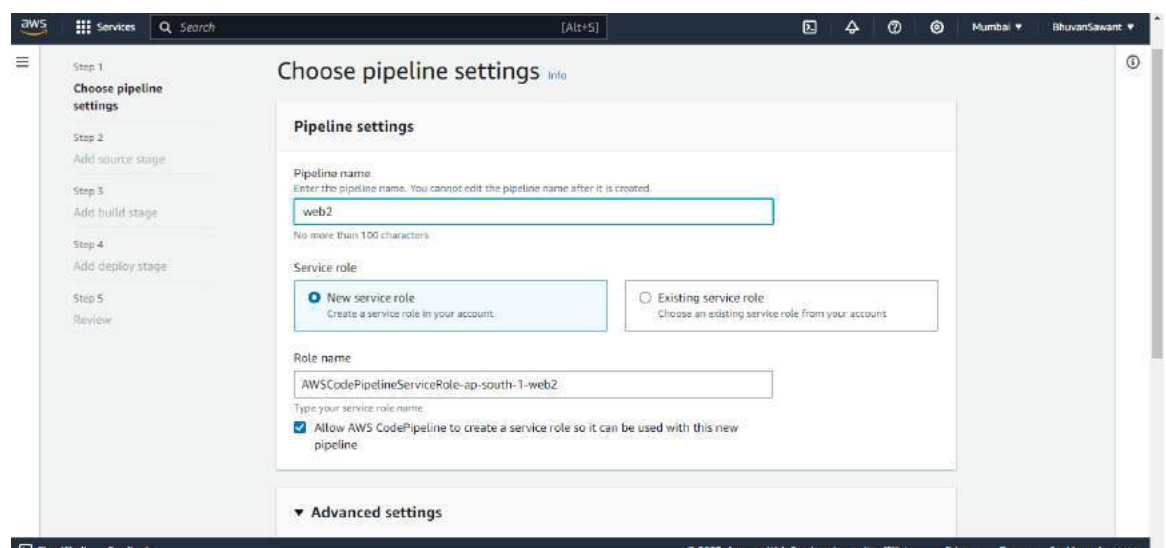
4. Artifact:

- Artifacts are the files and data that are produced as a result of an action. They are used to pass information between stages in a pipeline.

To deploy web application using CodePipeline here are the following steps to be followed:

1. Set Up Source Stage:

- Configure a source stage in AWS CodePipeline, linking to your version control system (e.g., CodeCommit, GitHub).



The screenshot shows the AWS CodePipeline console interface. On the left, a sidebar lists the steps of the pipeline: Step 1 (Choose pipeline settings), Step 2 (Add source stage), Step 3 (Add build stage), Step 4 (Add deploy stage), and Step 5 (Review). The main panel is titled 'Choose pipeline settings' and contains the following fields and options:

- Pipeline name:** A text input field containing 'web2'. Below it, a note states: 'Enter the pipeline name. You cannot edit the pipeline name after it is created. No more than 100 characters.'
- Service role:** Two radio button options: 'New service role' (selected) with the subtext 'Create a service role in your account.', and 'Existing service role' with the subtext 'Choose an existing service role from your account.'
- Role name:** A text input field containing 'AWSCodePipelineServiceRole-ap-south-1-web2'. Below it, a note states: 'Type your service role name.'
- Permissions:** A checked checkbox labeled 'Allow AWS CodePipeline to create a service role so it can be used with this new pipeline.'
- Advanced settings:** A section header with a downward arrow, currently collapsed.

2. Configure Build Stage:

- Set up a build stage using AWS CodeBuild to compile, test, and package your web application.

aws Services Search [Alt+S] Mumbai BhuvanSawant

Type your service role name

☒ Allow AWS CodePipeline to create a service role so it can be used with this new pipeline

▼ Advanced settings

Artifact store

☒ Default location
Use the default artifact store (Amazon S3 codepipeline-ap-south-1-36972723556) designated in the same region and account as your pipeline.

☐ Custom location
Choose an existing S3 location from your account in the same region and account as your pipeline.

Encryption key

☒ Default AWS Managed Key
Use the AWS managed customer master key for CodePipeline in your account to encrypt the data in the artifact store.

☐ Customer Managed Key
To encrypt the data in the artifact store under an AWS KMS customer managed key, specify the key ID, key ARN, or alias ARN.

Cancel Next

3. Define Deployment Stage:

- Create a deployment stage using AWS CodeDeploy or another deployment provider to deploy your application to target environments.

aws Services Search [Alt+S] Mumbai BhuvanSawant

Pipelines must have at least two stages. Your second stage must be either a build or deployment stage. Choose a provider for either the build stage or deployment stage.

Deploy

Deploy provider
Choose how you deploy to instances. Choose the provider, and then provide the configuration details for that provider.

Amazon S3

Region
Asia Pacific (Mumbai)

Bucket
mynews3bucketforpipeline

Deployment path - optional

☒ Extract file before deploy
The deployed artifact will be unzipped before deployment.

▼ Additional configuration

Next

4. Configure Approval (Optional):

- Optionally, add a manual approval stage to review and approve deployments before proceeding to the next stage.

Deployment path - optional

☒ Extract file before deploy
The deployed artifact will be unzipped before deployment.

▼ Additional configuration

KMS encryption key ARN - optional
Encrypt your object using a KMS encryption key. If no key is provided, objects remain un-encrypted.
arn:aws:kms:<region-ID>:<account-ID>:key/<key-ID> OR arn:aws:kms:<region-ID>:<account-ID>:alias/<alias-name>

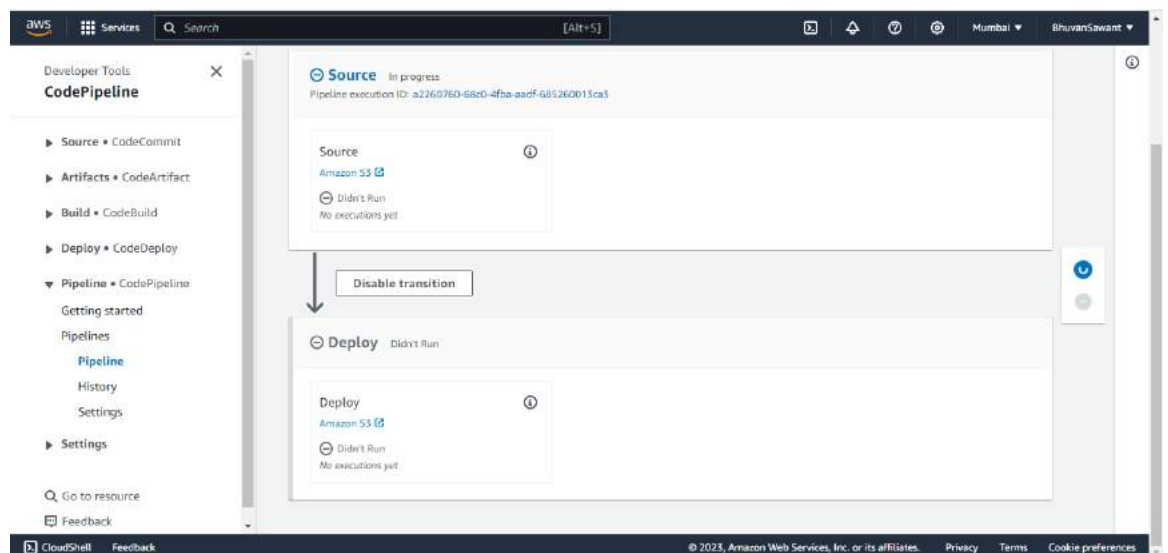
Canned ACL - optional
Specify an Amazon S3 canned access control list (ACL) for your bucket.
public-read

Cache control - optional
Set cache control for objects requested from your Amazon S3 bucket.

Cancel Previous Next

5. Artifact Passing:

- Ensure proper passing of artifacts between stages to maintain consistency in the deployment process.



6. Add Monitoring (Optional):

- Integrate monitoring tools (e.g., AWS CloudWatch) to track the performance and health of your application during and after deployment.

7. Configure Notifications (Optional):

- Set up notifications using AWS SNS or other services to receive alerts about pipeline events and status changes.

8. Test and Validate:

- Test the pipeline by triggering a build, ensuring that each stage executes successfully, and the application deploys as expected.

9. Modify Pipeline as Needed:

- Make adjustments to the pipeline configuration based on the specific requirements of your web application and deployment process.

10. Continuous Improvement:

- Implement continuous improvement practices, such as monitoring feedback, optimizing build and deployment scripts, and iterating on the pipeline structure.

CONCLUSION : In this assignment we learned how to host static web page through codepipeline and deploy web application using Code Pipeline.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 5

AIM : To understand the Kubernetes Cluster Architecture.

LO MAPPED : LO1,LO5

THEORY :

Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.

Let's take a look at why Kubernetes is so useful by going back in time.

- Traditional deployment era: Early on, organizations ran applications on physical servers. There was no way to define resource boundaries for applications in a physical server, and this caused resource allocation issues. For example, if multiple applications run on a physical server, there can be instances where one application would take up most of the resources, and as a result, the other applications would underperform. A solution for this would be to run each application on a different physical server. But this did not scale as resources were underutilized, and it was expensive for organizations to maintain many physical servers.
- Virtualized deployment era: As a solution, virtualization was introduced. It allows you to run multiple Virtual Machines (VMs) on a single physical server's CPU. Virtualization allows applications to be isolated between VMs and provides a level of security as the information of one application cannot be freely accessed by another application.
- Container deployment era: Containers are similar to VMs, but they have relaxed isolation properties to share the Operating System (OS) among the applications. Therefore, containers are considered lightweight. Similar to a VM, a container has its own filesystem, share of CPU, memory, process space, and more. As they are decoupled

from the underlying infrastructure, they are portable across clouds and OS distributions.

Why you need Kubernetes and what it can do

Kubernetes provides you with:

- Service discovery and load balancing Kubernetes can expose a container using the DNS name or using their own IP address. If traffic to a container is high, Kubernetes is able to load balance and distribute the network traffic so that the deployment is stable.
- Storage orchestration Kubernetes allows you to automatically mount a storage system of your choice, such as local storages, public cloud providers, and more.
- Automated rollouts and rollbacks You can describe the desired state for your deployed containers using Kubernetes, and it can change the actual state to the desired state at a controlled rate. For example, you can automate Kubernetes to create new containers for your deployment, remove existing containers and adopt all their resources to the new container.
- Automatic bin packing You provide Kubernetes with a cluster of nodes that it can use to run containerized tasks. You tell Kubernetes how much CPU and memory (RAM) each container needs. Kubernetes can fit containers onto your nodes to make the best use of your resources.
- Self-healing Kubernetes restarts containers that fail, replaces containers, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.
- Secret and configuration management Kubernetes lets you store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys. You can deploy and update secrets and application configuration without rebuilding your container images, and without exposing secrets in your stack configuration.

Q1)What are the various Kubernetes services running on nodes? Describe the role of each service.

Here are the key services running on Kubernetes nodes and their roles:

- **Container Runtime:** Docker (or another container runtime): The container runtime is responsible for running and managing containers. Docker is a common choice, but Kubernetes can also work with container runtimes like containerd and CRI-O. Its role is to create, start, stop, and manage containers based on container images.
- **Kubelet:** The Kubelet is an agent running on each node that communicates with the Kubernetes control plane (master node). It ensures that containers are running in a Pod (the smallest deployable unit in Kubernetes). Kubelet pulls container images and starts, stops, and restarts containers as needed.
- **Kube Proxy:** Kube Proxy is responsible for network connectivity within the cluster. It maintains network rules on each node to enable communication between different Pods. It sets up and manages network routes, ensuring that services and Pods can connect to each other.
- **Container Network Interface (CNI):** CNI plugins provide the network abstraction needed for connecting containers within a Pod and across Pods on different nodes. Kubernetes allows various CNI plugins like Calico, Flannel, and Weave, which implement different network policies and features.
- **Node Status:** The Node Status service periodically reports the node's condition to the control plane. This information helps the control plane make scheduling decisions. It communicates the node's available resources, capacity, and overall health.

Q2) What is Pod Disruption Budget (PDB)?

A Pod Disruption Budget (PDB) is a policy object in Kubernetes that allows you to define and control the allowable disruption or unavailability of a set of pods during voluntary disruptions, such as node draining or maintenance activities. PDBs are used to ensure that critical workloads have a certain level of availability and are not unduly disrupted when nodes need to be updated, scaled down, or otherwise modified.

Key aspects of a Pod Disruption Budget (PDB) include:

- **Targeted Pods:** PDBs are applied to a specific set of pods by specifying a set of labels and selectors. This defines which pods the PDB applies to.
- **MaxUnavailable:** The most crucial aspect of a PDB is specifying the maxUnavailable field. This field determines the maximum number of pods that can be unavailable during a disruption event. It can be specified as an absolute number (e.g., 1) or a percentage (e.g., 10%). This ensures that a minimum number of pods within the defined set remain available.
- **Selector and Labels:** The PDB uses selector rules to select pods based on their labels. Only pods matching the selector rules will be subject to the disruption budget.
- **Availability Constraints:** PDBs can be used to enforce constraints on the availability of pods, helping to maintain a desired level of redundancy and fault tolerance for important workloads.
- **Preventing Disruptions:** When the disruption budget is violated, Kubernetes will prevent any action that would cause more pods to become unavailable than specified by the PDB. This can include scaling down replicas, draining nodes, or terminating pods.

Q3) What is the role of Load Balance in Kubernetes?

Load balancing in Kubernetes plays a critical role in ensuring the availability, scalability, and reliability of applications and services running within the cluster. It involves distributing incoming network traffic or requests across multiple pods or replicas of an application, thereby optimizing resource utilization and improving fault tolerance. Here are the key roles and benefits of load balancing in Kubernetes:

- **Distribute Traffic:** Load balancers evenly distribute incoming traffic or requests across a set of backend pods or services. This distribution prevents overloading individual pods and ensures efficient utilization of resources.
- **High Availability:** Load balancing helps ensure high availability of services by distributing traffic across multiple pods. In the event of a

pod failure, the load balancer redirects traffic to healthy pods, minimizing service disruption.

- **Scaling:** As your application traffic increases, you can scale your application horizontally by adding more pods or replicas. Load balancers automatically adapt to the increased capacity, spreading traffic across the expanded set of pods.
- **Service Discovery:** Load balancers provide a stable endpoint (e.g., a virtual IP address) for clients to access a service. The actual pods behind the service can change (e.g., due to scaling, updates, or failures), but clients don't need to be aware of these changes.
- **Health Checking:** Load balancers can perform health checks on backend pods to verify their availability and responsiveness. If a pod becomes unhealthy, the load balancer stops sending traffic to it, ensuring that clients only connect to healthy instances.
-

CONCLUSION : Understood about Kubernetes and Kubernetes cluster architecture

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 6

AIM : To understand terraform lifecycle, core concepts/ terminologies and install it on a linux machine.

LO MAPPED : LO1,LO5

THEORY :

Terraform, developed by HashiCorp, is an open-source infrastructure as code (IaC) software tool that allows users to define and provision data center infrastructure using a declarative configuration language. In simpler terms, it lets you codify your infrastructure, enabling consistent and reproducible deployments. It is used for defining, provisioning, and managing cloud infrastructure using a declarative language.

Core Concepts:

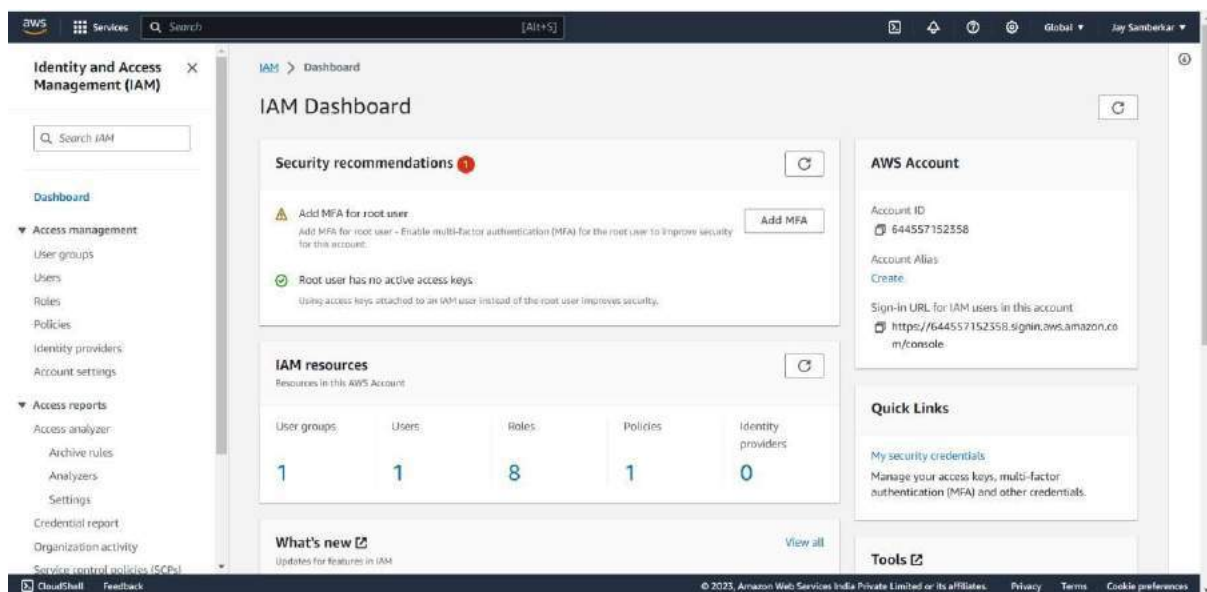
- **Providers:** Terraform uses providers to interact with cloud services. Examples include AWS, Azure, Google Cloud, and many others. Each provider offers resource types that can be managed.
- **Resources:** These are the primary components in Terraform. A resource might be a physical component such as an EC2 instance in AWS or a database in Azure.
- **State:** Terraform maintains a state file that maps real-world resources to your configuration. This state is used to determine what Terraform will do on the next apply.
- **Modules:** These are containers for multiple resources that are used together. They provide a way to group resources, and they can be used to create reusable infrastructure components.
- **Variables and Outputs:** Variables allow for parameterization of the Terraform configuration, making it more dynamic and flexible.

Outputs are a way to get information about the infrastructure, like IP addresses or DNS names.

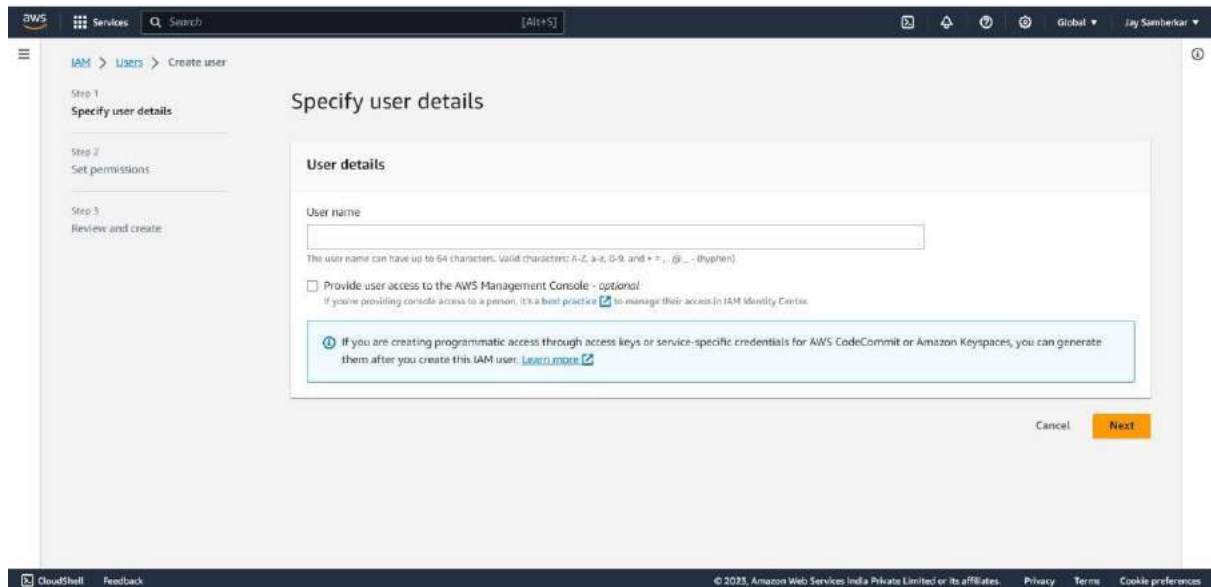
- Provisioners: While not always recommended due to their imperative nature, provisioners can be used to model specific actions on the local machine or on a remote machine, like executing scripts.

Steps :

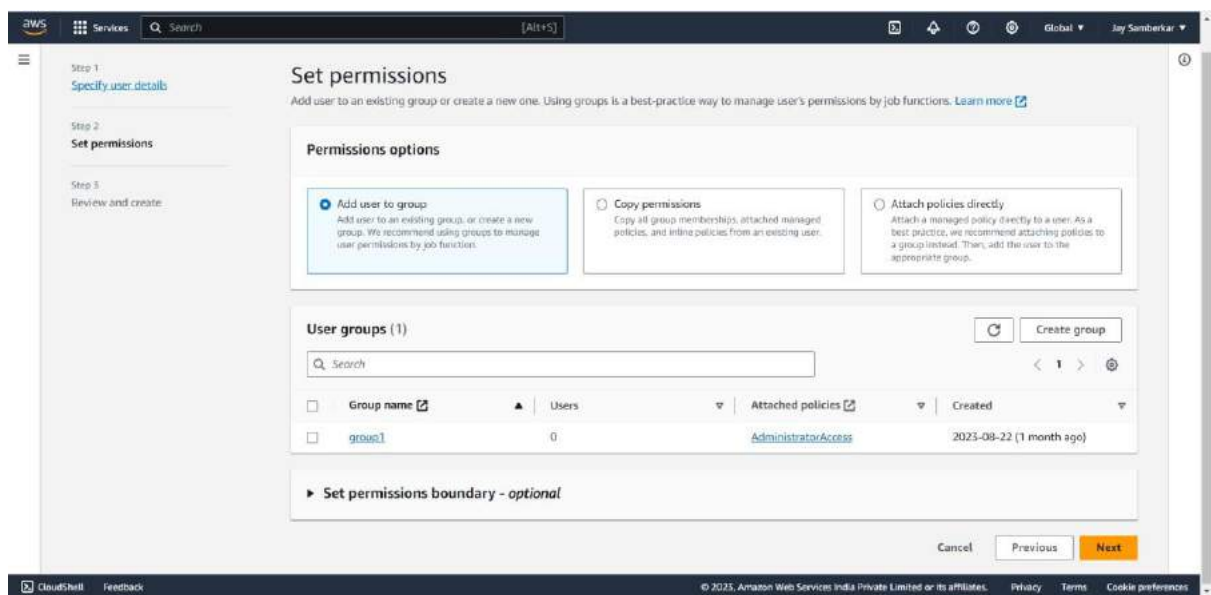
1. Login and Search IAM on AWS Console and select the first option.



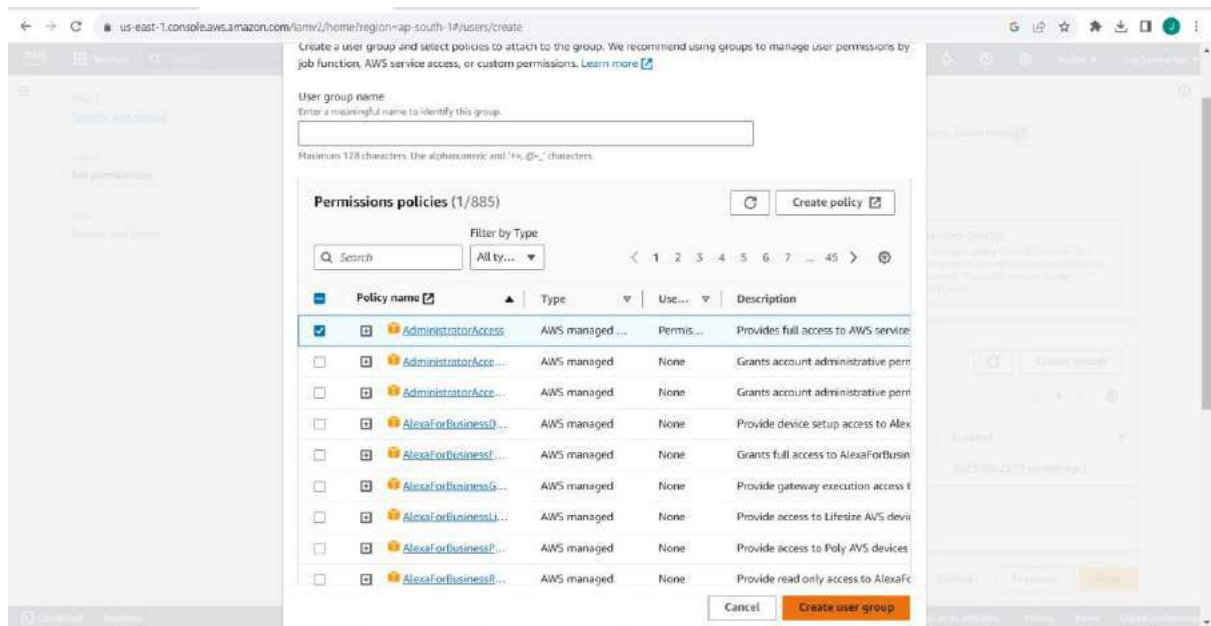
2. Head to users and Click on Create User



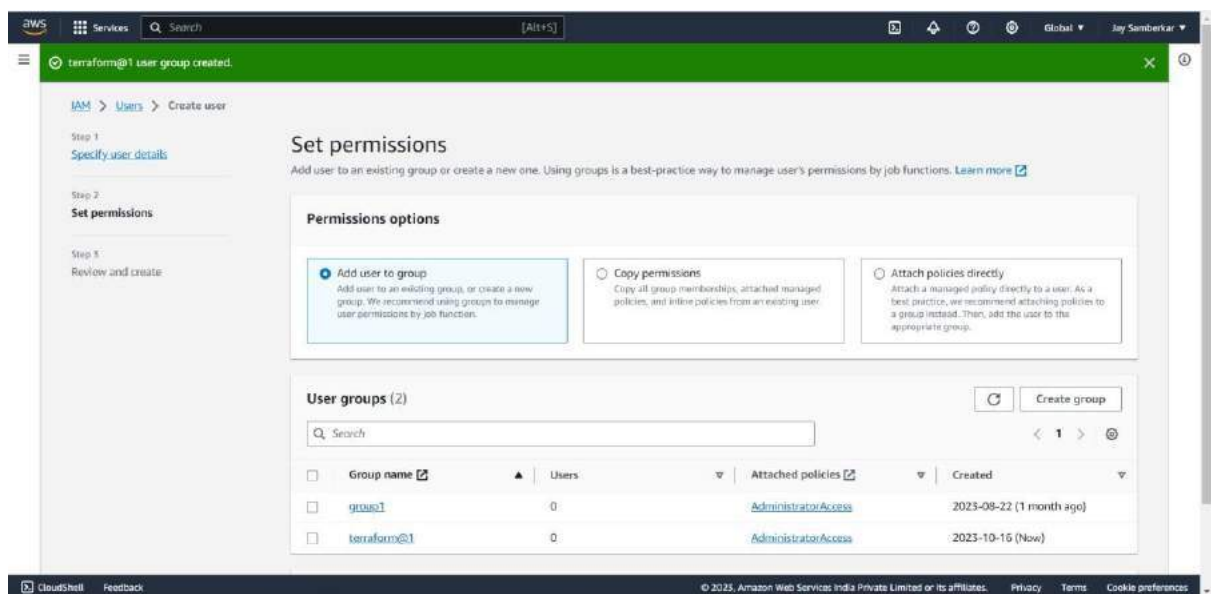
3. Enter a new username and click on next
4. Click on create group



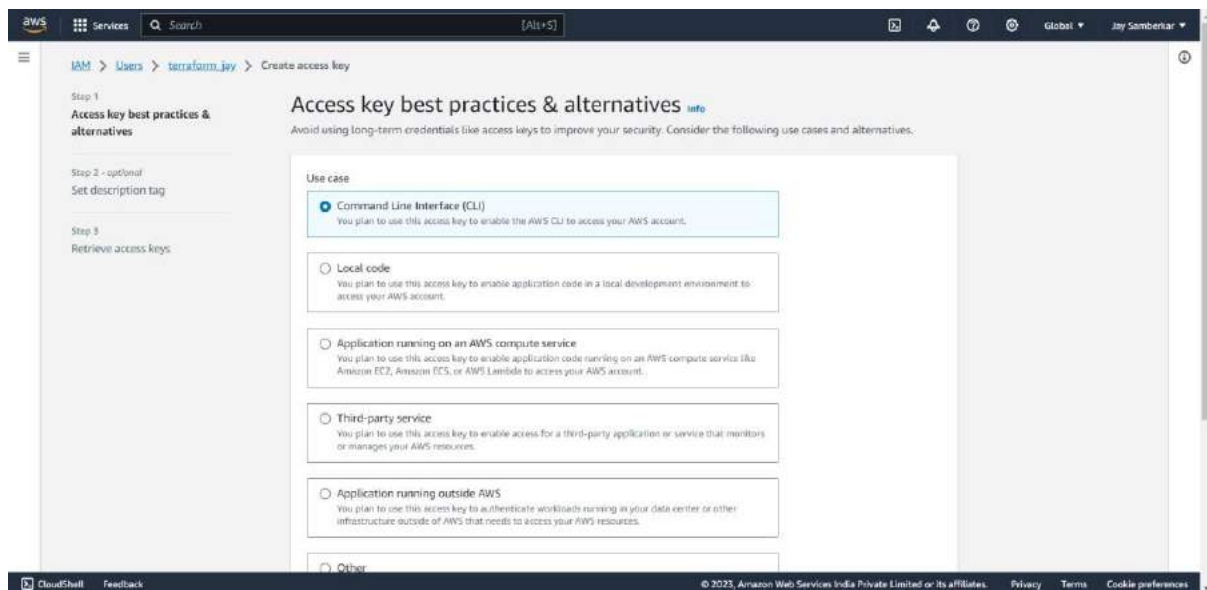
5. Select the first Policy will full access and enter a group name and create the group.



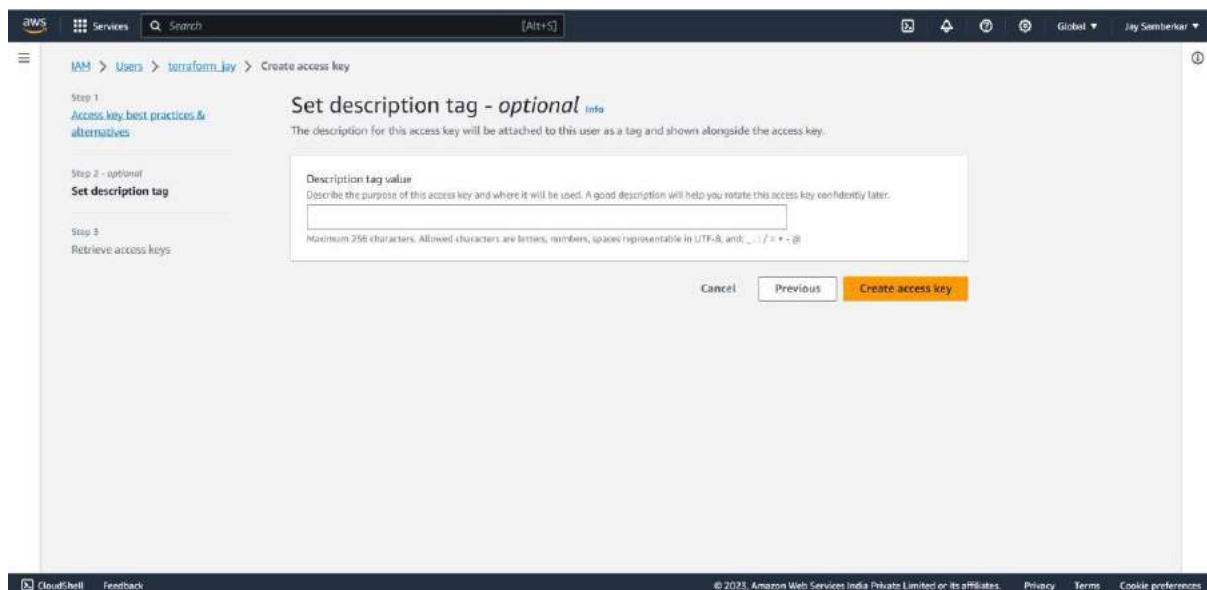
6. Then click on next, select the policy and create user.



7. Then create an access key with CLI.



8. Leave the description field empty and create a key. Download the access key in the .csv file.



9. Install terraform.


```

Last login: Sat Aug 19 17:44:56 on ttys000
> brew tap hashicorp/tap
Running `brew update --auto-update`...
Installing from the API is now the default behaviour!
You can save space and time by running:
  brew untap homebrew/core
  brew untap homebrew/cask
==> Auto-updated Homebrew!
Updated 3 taps (mongodb/brew, homebrew/core and homebrew/cask).
==> New Formulae
arm-none-eabi-binutils      medusa
arm-none-eabi-gcc           mjml
arm-none-eabi-gdb           mongodb/brew/mongodb-community@6.0
asnmap                     mongodb/brew/mongodb-enterprise@6.0
cargo-auditable            mongodb/brew/mongodb-mongocryptd@6.0
cdi                        mysql-client@8.0
cloudlist                 mysql@8.0
coder                     ollama
ctpv                     proxify
czkawka                  python-certifi
dnsrobocert              riff
dolphie                  riscv64-elf-binutils
ebook2cw                 riscv64-elf-gcc
go@1.20                  riscv64-elf-gdb
img2pdf                  rpmspectool

```

10. Create the terraform config file main.tf with required configurations.

```

> nano main.tf
> terraform init

Initializing the backend...

Initializing provider plugins...
- Finding hashicorp/aws versions matching "~> 4.16"...
- Installing hashicorp/aws v4.67.0...
- Installed hashicorp/aws v4.67.0 (signed by HashiCorp)

Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
~/terraform_scripts 25s

```

11. Run the command terraform plan and enter yes

If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.

```
> terraform plan
```

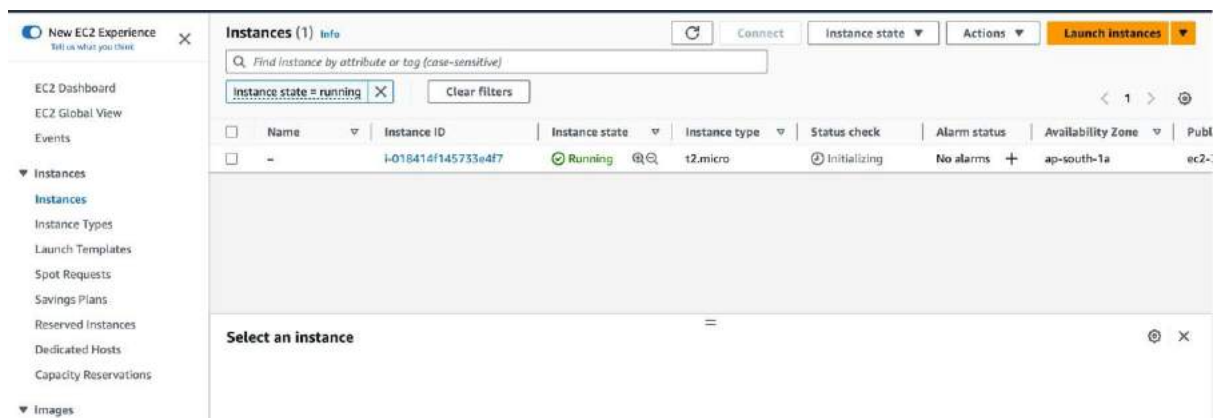
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# aws_instance.Ubuntu will be created
+ resource "aws_instance" "Ubuntu" {
  + ami                  = "ami-0f5ee92e2d63afc18"
  + arn                  = (known after apply)
  + associate_public_ip_address = (known after apply)
  + availability_zone     = (known after apply)
  + cpu_core_count       = (known after apply)
  + cpu_threads_per_core = (known after apply)
  + disable_api_stop      = (known after apply)
  + disable_api_termination = (known after apply)
  + ebs_optimized         = (known after apply)
  + get_password_data     = false
  + host_id               = (known after apply)
  + host_resource_group_arn = (known after apply)
```

12. Check EC2 instances on AWS, you'll find an instance running started by terraform.



13. Now run terraform destroy to terminate the instance.

```

- enable_resource_name_dns_a_record    = false -> null
- enable_resource_name_dns_aaaa_record = false -> null
- hostname_type                        = "ip-name" -> null
}

- root_block_device {
  - delete_on_termination = true -> null
  - device_name           = "/dev/sda1" -> null
  - encrypted             = false -> null
  - iops                  = 100 -> null
  - tags                 = {} -> null
  - throughput            = 0 -> null
  - volume_id             = "vol-07a64df7d00e4e7a1" -> null
  - volume_size           = 8 -> null
  - volume_type           = "gp2" -> null
}
}

```

Plan: 0 to add, 0 to change, 1 to destroy.

Do you really want to destroy all resources?

Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

CONCLUSION : We understood the workings of Terraform to create and manage EC2 instances in AWS which provides a highly automated and reproducible infrastructure-as-code solution. The ability to effortlessly spin up and tear down instances not only enhances operational agility but also ensures optimal resource utilization.

NAME : SOUMIL SALVI

ROLL NO : 104

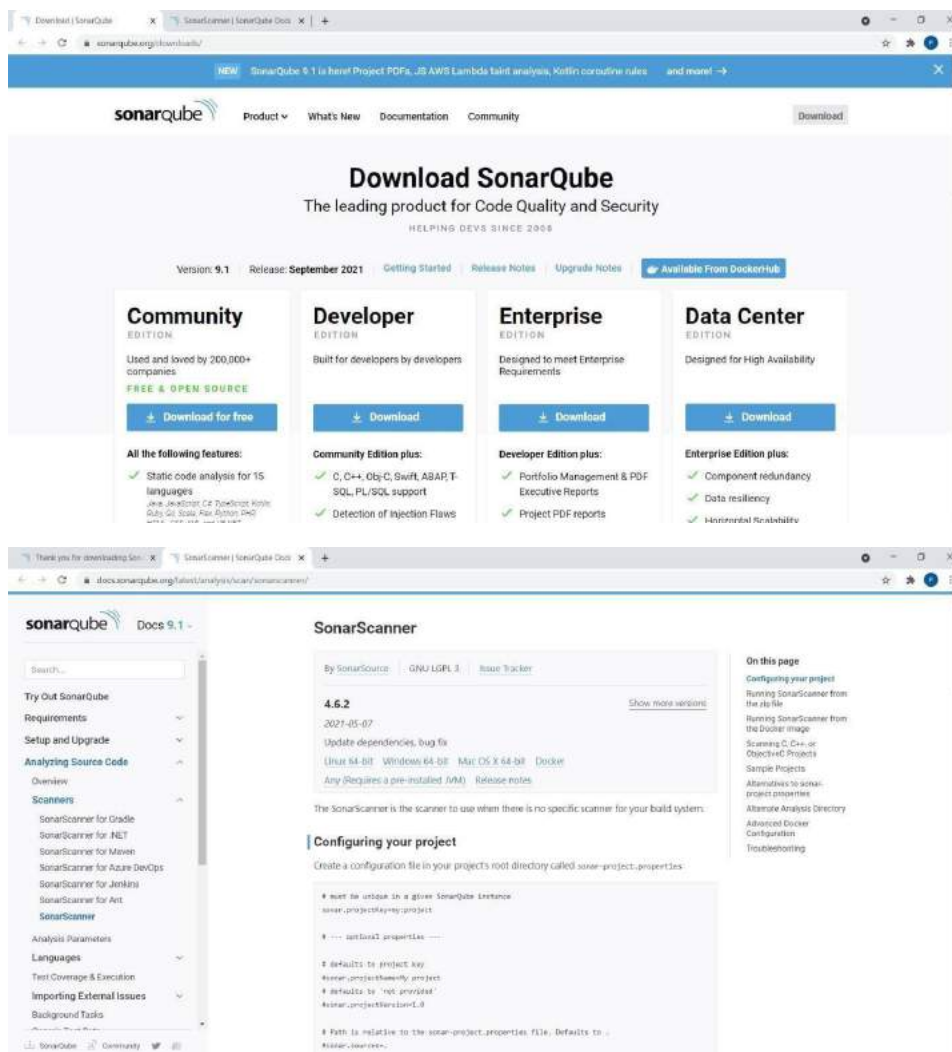
ASSIGNMENT 7

AIM : To perform static analysis on python programs using SonarCube SAST processes

LO MAPPED : LO4

THEORY :

Download SonarQube and Sonar Scanner



After downloading, set Environment Variables. Add “sonarqube-9.1.0.47736\bin” to Path.

- cd "sonarqube-9.1.0.47736\bin\windows-x86-64"
- StartSonar.bat

[illegible]

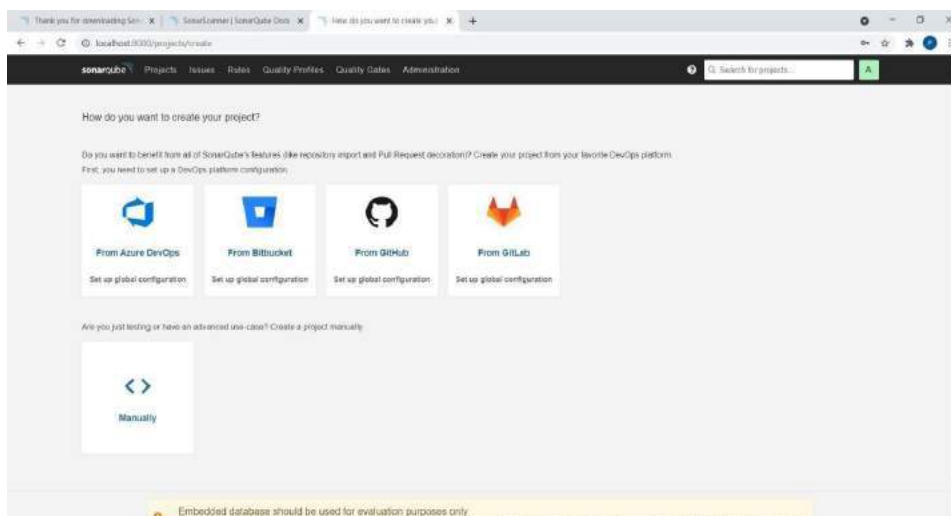
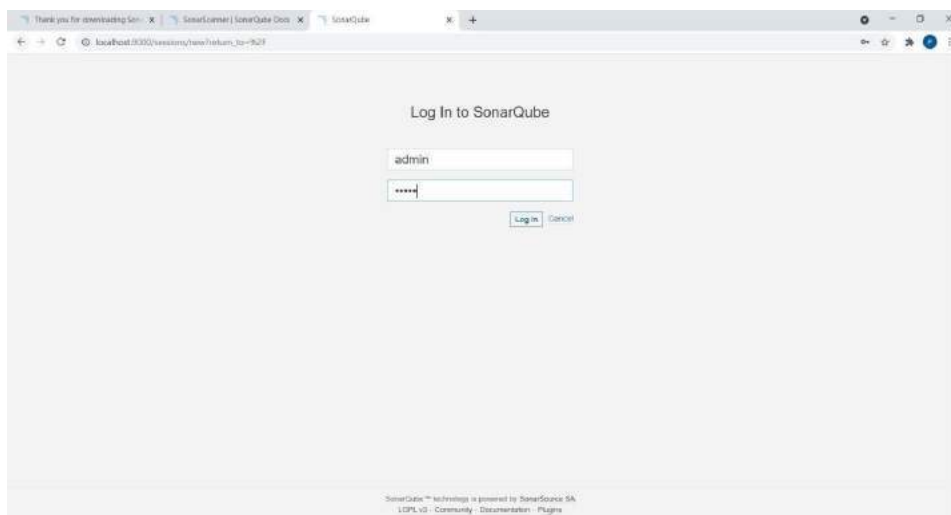
- cd "sonar-scanner-4.6.2.2472-windows\bin"
- sonar-scanner


```
Command Prompt
C:\Users\Priyansi\Downloads\sonar-scanner-4.6.2.2472-windows\bin>sonar-scanner
INFO: Scanner configuration file: C:\Users\Priyansi\Downloads\sonar-scanner-4.6.2.2472-windows\bin\..\conf\sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: SonarScanner 4.6.2.2472
INFO: Java 11.0.11 AdoptOpenJDK (64-bit)
INFO: Windows 10 10.0 amd64
INFO: User cache: C:\Users\Priyansi\.sonar\cache
INFO: Scanner configuration file: C:\Users\Priyansi\Downloads\sonar-scanner-4.6.2.2472-windows\bin\..\conf\sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: Analyzing on SonarQube server 9.1.0
INFO: Default locale: "en_IN", source code encoding: "windows-1252" (analysis is platform dependent)
INFO: Load global settings
INFO: -----
INFO: EXECUTION FAILURE
INFO: -----
INFO: Total time: 3.058s
INFO: Final Memory: 5M/20M
INFO: -----
ERROR: Error during SonarScanner execution
ERROR: Not authorized. Analyzing this project requires authentication. Please provide a user token in sonar.login or other credentials in sonar.login and sonar.password.
ERROR:
ERROR: Re-run SonarScanner using the -X switch to enable full debug logging.

C:\Users\Priyansi\Downloads\sonar-scanner-4.6.2.2472-windows\bin>
```

Server up and running on localhost:9000

Login using credentials as User: admin and Password: admin and Set a new password



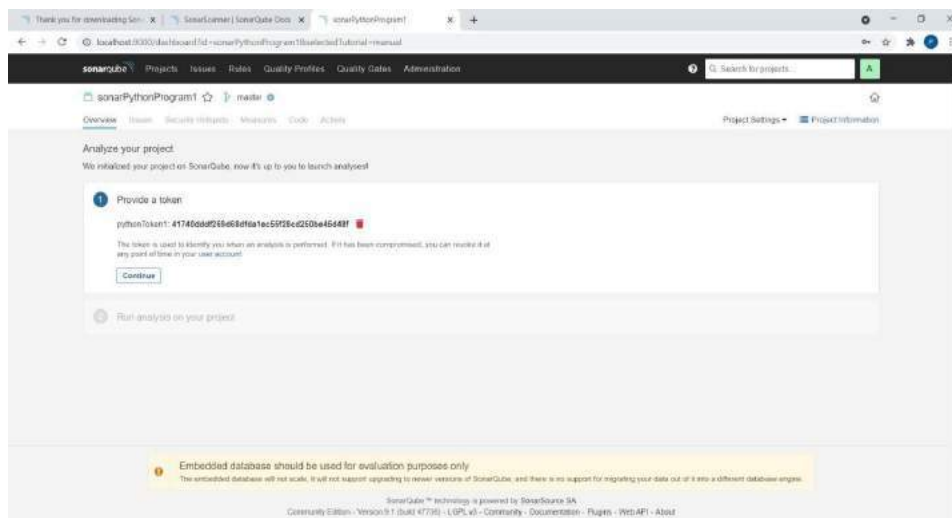
Click on Create a project Manually

Give any Project display name.

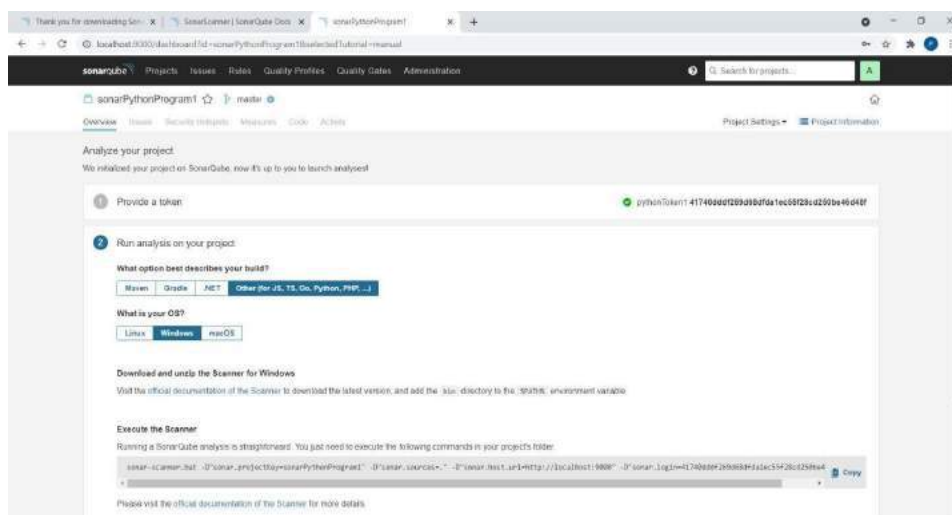
Click on Locally

The screenshot shows the SonarQube web interface. At the top, there's a navigation bar with links like 'sonarQube', 'Projects', 'Issues', 'Rules', 'Quality Profiles', 'Quality Gates', and 'Administration'. Below this, the breadcrumb trail shows 'sonarPythonProgram1' and 'main'. The main heading is 'Analyze your project'. A sub-heading states: 'We initialized your project on SonarQube, now it's up to you to launch analysis!'. The first step, 'Provide a token', is highlighted. It contains a 'Generate a token' button, a text input field with the placeholder 'pythonToken1', and a 'Generate' button. A note explains that the token is used for authentication and can be recovered if compromised. The second step, 'Run analysis on your project', is partially visible below.

Give any name to token and click on Generate.



Click on Continue.



Save a Python program in a folder.

```
class Solution(object):
```

```
def romanToInt(self, s):
```

```
    roman =
```

```
    {'I':1,'V':5,'X':10,'L':50,'C':100,'D':500,'M':1000,'IV':4,'IX':9,'XL':40,'XC':90,'CD':400,'CM':900}
    i = 0
    num = ""
    while i < len(s):
```

```
        if i+1<len(s) and s[i:i+2] in roman:
```

```
            num+=roman[s[i:i+2]]
            i+=2
```

```
        else:
            #print(i)
```

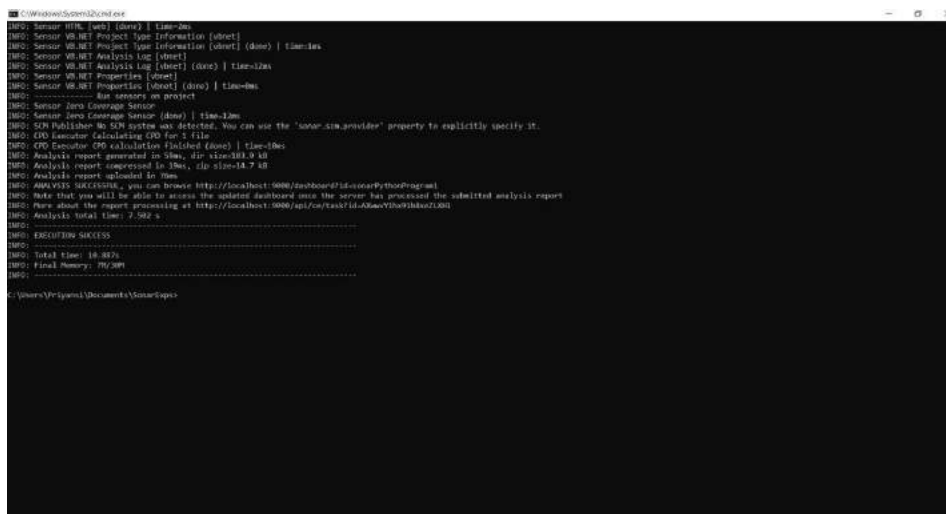


```
num+=roman[s[i]] i+=1
```

```
return num
```

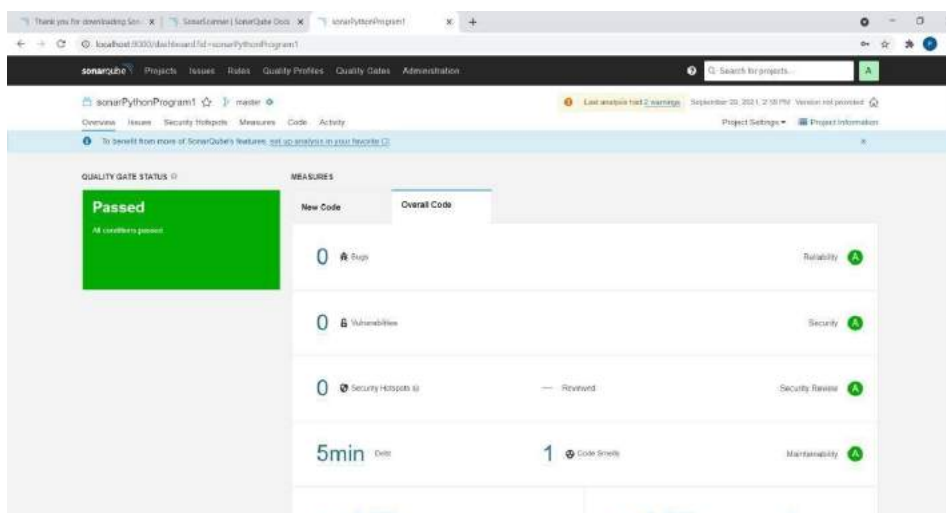
```
ob1 = Solution() print(ob1.romanToInt("III")) print(ob1.romanToInt("CDXLIII"))
```

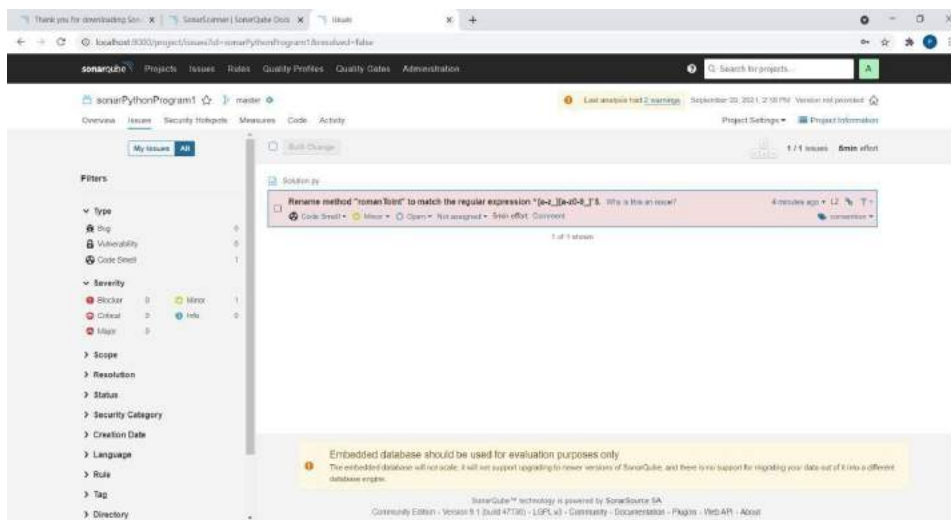
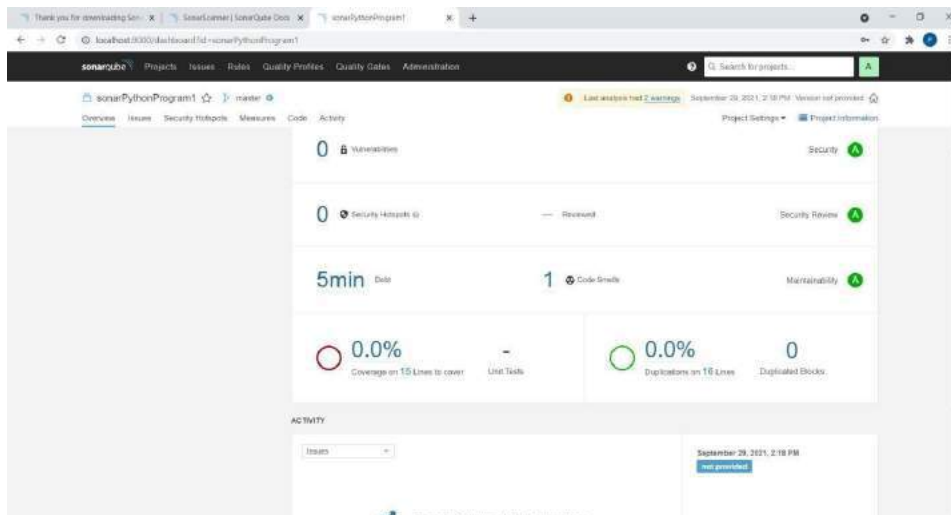
Open command prompt in this folder and Run program using copied command.
sonar-scanner.bat -D"sonar.projectKey=<YourDisplayName>"
-D"sonar.sources=." -D"sonar.host.url=http://localhost:9000" –



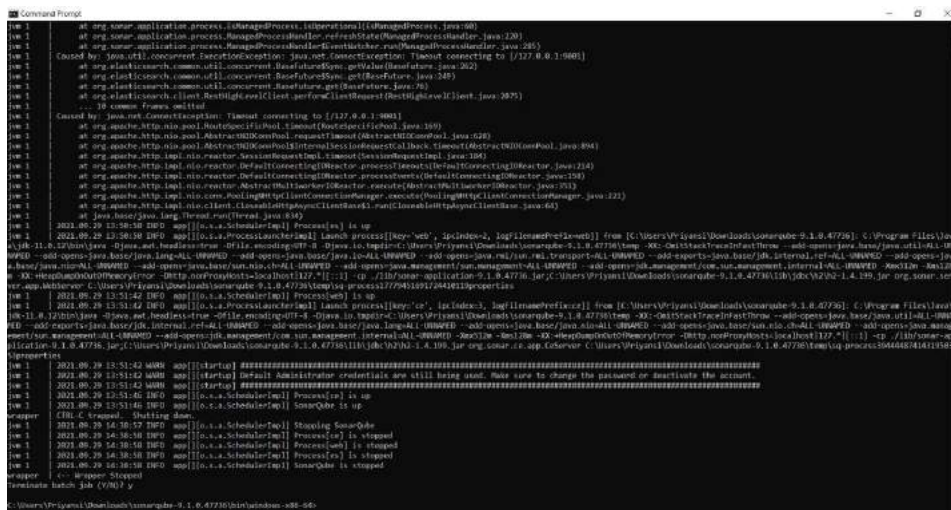
```
C:\Windows\system32\cmd.exe
DMS: Sensor VB.NET [web] [done] | time=1ms
DMS: Sensor VB.NET Project Type Information [direct] | time=1ms
DMS: Sensor VB.NET Project Type Information [direct] [done] | time=1ms
DMS: Sensor VB.NET Analysis Log [direct]
DMS: Sensor VB.NET Analysis Log [direct] [done] | time=12ms
DMS: Sensor VB.NET Properties [direct]
DMS: Sensor VB.NET Properties [direct] [done] | time=1ms
DMS: Sensor Java Coverage Sensor
DMS: Sensor Java Coverage Sensor [done] | time=11ms
DMS: Sensor SCM Publisher No SCM system was detected. You can use the 'sonar.scm.provider' property to explicitly specify it.
DMS: CPO Executor calculating CPO for 1 file
DMS: CPO Executor CPO calculation finished [done] | time=10ms
DMS: Analysis report generated in 5ms, dir size=101.9 KB
DMS: Analysis report compressed in 5ms, dir size=10.7 KB
DMS: Analysis report uploaded in 10ms
DMS: ANALYSIS SUCCESSFUL, you can browse http://localhost:9000/dashboard?id=sonarPythonProgram1
DMS: Here about the report processing at http://localhost:9000/api/cx/task?id=0ba0c70a91b0a71001
DMS: Analysis total time: 7.562 s
DMS: EXECUTION SUCCESS
DMS: Total time: 10.807s
DMS: Final Memory: 796,389
DMS:
C:\Users\jpr\Documents\sonar\sonar-scanner\bin\windows\sonar-scanner.bat
```

Given below is the inspection of code quality to perform automatic reviews with static analysis of code to detect bugs, code smells, and security vulnerabilities.





Press “Ctrl + C” to stop the server.



CONCLUSION : In this assignment we implemented static analysis on python programs using SonarCube SAST processes.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 8

AIM : To understand continuous monitoring and Installation and configuration of Nagios on Linux Machine

LO MAPPED : LO1,LO5

THEORY :

Login to Aws.

Create Ec2 instances on Aws account of any linux os

Then Run the following command in SS

```
ubuntu@ip-172-31-13-219:~$ sudo apt update
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu focal InRelease
Get:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Get:3 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
Get:4 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu focal-backports InRelease [101 kB]
Get:5 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu focal/universe amd64 Packages [8628 kB]
Get:6 http://security.ubuntu.com/ubuntu focal-security/main amd64 Packages [909 kB]
Get:7 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu focal/universe Translation-en [5124 kB]
```

`cd /usr/src/`

`sudo wget https://github.com/NagiosEnterprises/nagioscore/archive/nagios-4.4.2.tar.gz`

```
ubuntu@ip-172-31-13-219:/usr/src$ sudo wget https://github.com/NagiosEnterprises/nagioscore/archive/nagios-4.4.2.tar.gz
--2021-10-05 10:24:05-- https://github.com/NagiosEnterprises/nagioscore/archive/nagios-4.4.2.tar.gz
Resolving github.com (github.com)... 13.234.176.102
Connecting to github.com (github.com)|13.234.176.102|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://codeload.github.com/NagiosEnterprises/nagioscore/tar.gz/nagios-4.4.2 [following]
--2021-10-05 10:24:05-- https://codeload.github.com/NagiosEnterprises/nagioscore/tar.gz/nagios-4.4.2
Resolving codeload.github.com (codeload.github.com)... 13.233.43.20
Connecting to codeload.github.com (codeload.github.com)|13.233.43.20|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified [application/x-gzip]
Saving to: 'nagios-4.4.2.tar.gz'

nagios-4.4.2.tar.gz           [ 10.78M  16.9MB/s  in 0.6s
2021-10-05 10:24:06 (16.9 MB/s) - 'nagios-4.4.2.tar.gz' saved [11301457]
```

`sudo tar xzf nagios-*.tar.gz cd`

`nagioscore-nagios-*/`

```
ubuntu@ip-172-31-13-219:/usr/src$ cd nagioscore-nagios-*/
ubuntu@ip-172-31-13-219:/usr/src/nagioscore-nagios-4.4.2$
```

Now finally run the following command

```
sudo ./configure --with-httpd-conf=/etc/apache2/sites-enabled
```

if error comes install c compiler on the linux by
following this link

<https://linuxize.com/post/how-to-install-gcc-compiler-on-ubuntu-18-04/> Finally-

```
4. Rerun the configure script.

NOTE: If you can't get the configure script to recognize the GD libs
on your system, get over it and move on to other things. The
CGIs that use the GD libs are just a small part of the entire
Nagios package. Get everything else working first and then
revisit the problem. Make sure to check the nagios-users
mailing list archives for possible solutions to GD library
problems when you resume your troubleshooting.

*****

checking ltdl.h usability... no
checking ltdl.h presence... no
checking for ltdl.h... no
checking dlfcn.h usability... yes
checking dlfcn.h presence... yes
checking for dlfcn.h... yes
checking for dlopen in -ldl... yes
checking for extra flags needed to export symbols... -Wl,-export-dynamic
checking for linker flags for loadable modules... -shared
checking for traceroute... no
checking for type va_list... yes
checking for perl... /usr/bin/perl
checking for unzip... no
```

Now let us install plugins by

```
sudo wget -O nagios-plugins.tar.gz https://github.com/nagios-plugins/nagiosplugins/archive/release-2.2.1.tar.gz
```

then

```
sudo tar xzf nagios-plugins.tar.gz then
```

```
cd nagios-plugins-release-2.2.1
```

finally start and then check status of nagi os

sudo systemctl start nagios sudo

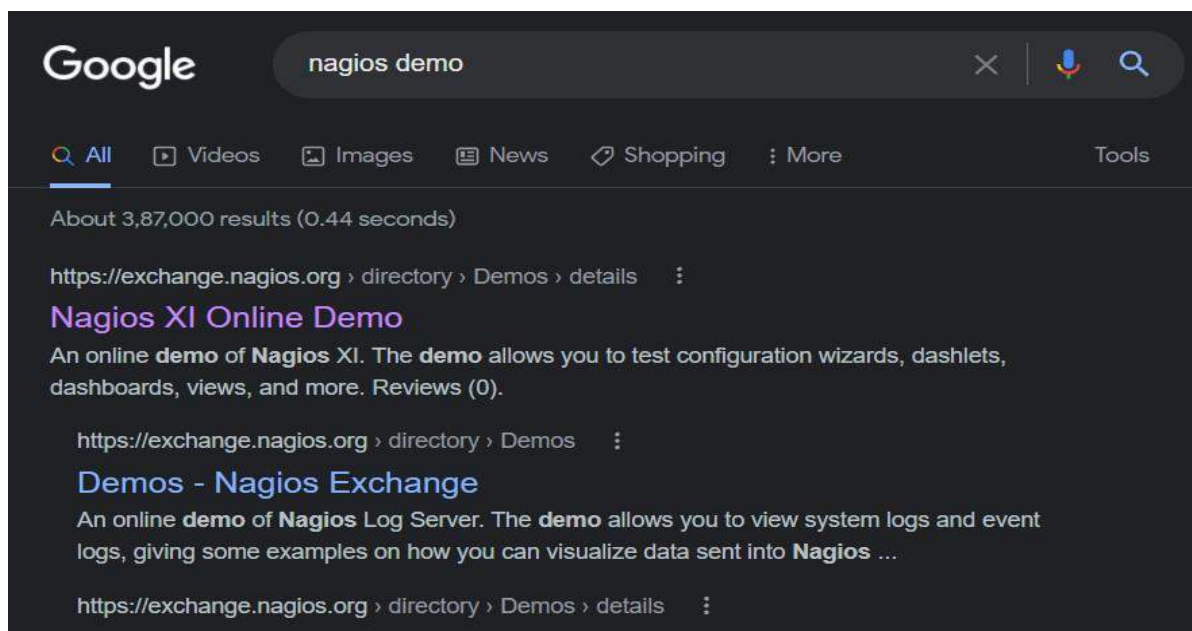
systemctl status nagios

```
* nagios.service - Nagios Core 4.4.2
   Loaded: loaded (/lib/systemd/system/nagios.service; enabled; vendor preset: enable
   Active: active (running) since Fri 2018-11-16 14:54:21 PST; 1s ago
     Docs: https://www.nagios.org/documentation
   Process: 18294 ExecStopPost=/bin/rm -f /usr/local/nagios/var/rw/nagios.cmd (code=ex
   Process: 18293 ExecStop=/bin/kill -s TERM ${MAINPID} (code=exited, status=0/SUCCESS
   Process: 18315 ExecStart=/usr/local/nagios/bin/nagios -d /usr/local/nagios/etc/nag
   Process: 18313 ExecStartPre=/usr/local/nagios/bin/nagios -v /usr/local/nagios/etc/
 Main PID: 18325 (nagios)
    Tasks: 6 (limit: 2319)
   CGroup: /system.slice/nagios.service
```

Steps-

Go to google.com, Search NagiOs Demo

Click on the first link shown below



Now click on the website-

Nagios®

[Network:](#) | [Enterprise](#) | [Support](#) | [Library](#) | [Project](#) | [Exchan](#)

[Home](#) | [Directory](#) | [About](#)

[Home](#) | [Directory](#) | [Demos](#) | [Nagios XI Online Demo](#)

Directory Tree

Nagios XI Online Demo

[Submit review](#) | [Recommend](#) | [Print](#) | [Visit](#) | [Claim](#)

Rating

Favoured: 0

0 votes

Owner **egalstad**

Website **nagiosxi.demos.nagios.com**

Hits 141800

Search Exchange

[Advanced Search](#)

Search All Sites

Nagios Live Webinars

Let our experts show you how Nagios can help your organization.

Now click on login as administrator

[Nagios XI](#)
[Login](#)

Login

[Forgot your password?](#)

Select Languages:

Nagios XI Demo System

Demo Account Options

You can access the demo with different accounts to get a different view of the monitoring system.

- Administrator Access** - An account with administrative privileges.

Username: nagiosadmin
Password: nagiosadmin
- Read-Only User Access** - A user account that can view all hosts and services, but not make any changes or issue commands.

Username: readonly
Password: readonly
- Advanced User Access** - An advanced user account that can see and control (schedule downtime, edit, etc.) all hosts and services.

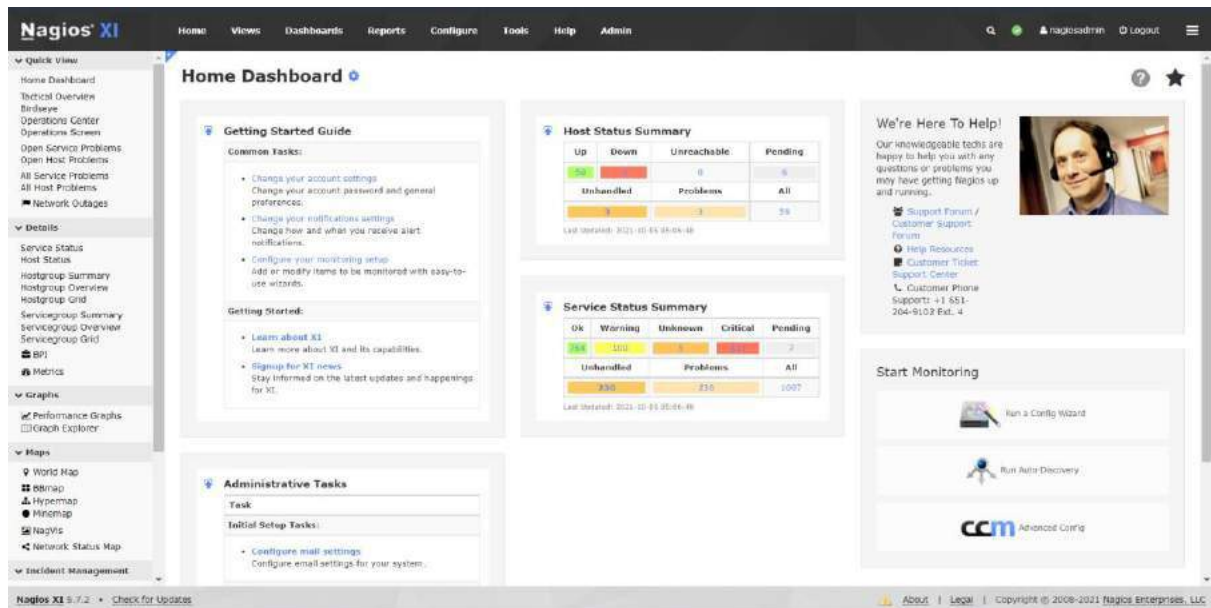
Username: advanced
Password: advanced
- Normal User Access** - A sample "normal" user account that has rights to view only a defined subset of all hosts and services.

Username: jdoe
Password: jdoe
- Administrator Access** - showing the dark theme.

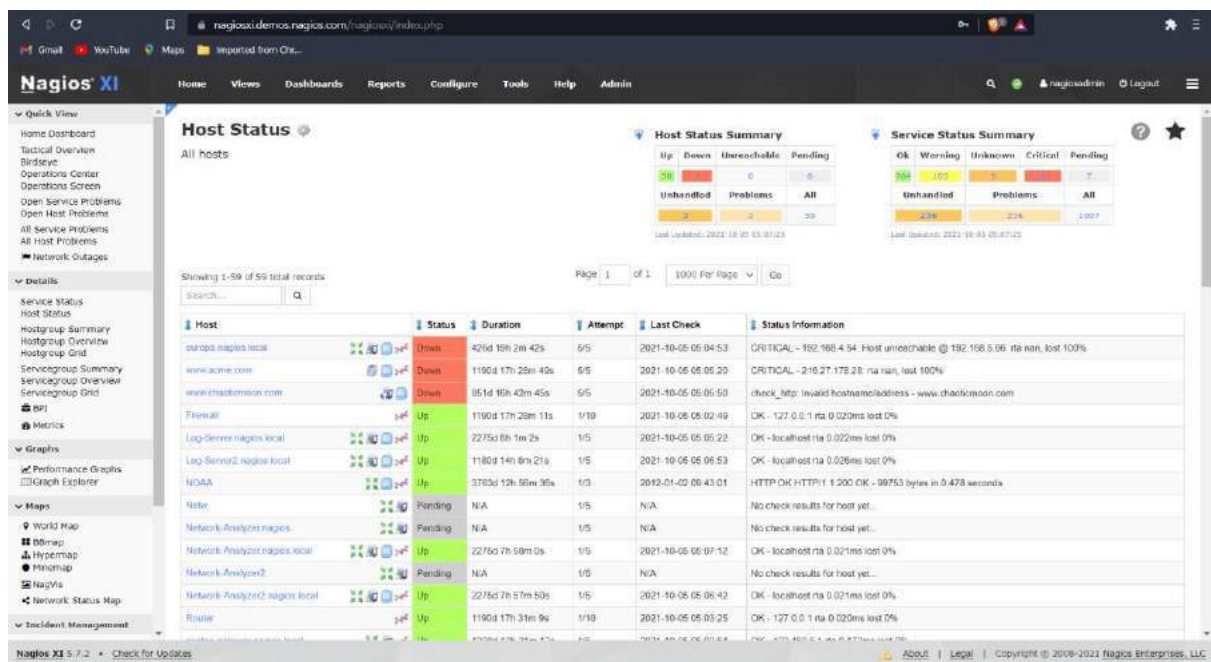
Username: darktheme
Password: darktheme

Demo Notes

It will have interface like this



Now click on Host status-



In the above image one can see Host Status Summary and Service Status Summary also how many host are up, down and also errors in detail Now click on Host Group Status.

✓ Disk Usage
404 Errors
apache
Apache 404 Errors
Apache Web Server
Aurora Activity
Bandwidth Spike
CPU State
CPU Usage
Cron Scheduling Daemon
Current Load
Current Users
DNS IP Match
DNS IP Match - www.acme.c
DNS Resolution
DNS Resolution - www.acme
Dovecot Mail Server
Drive C: Disk Usage
Explorer
Facebook Usage
Failed SSH Logins
HTTP
IIS Web Server
Linux Failed Logins
Load
Login Errors
MySQL Active Transactions
MySQL Average Wait Time
MySQL Buffer Hit Ratio
MySQL Checkpoint Pages Pe
MySQL Connection Time
MySQL Database Pages
MySQL Deadlocks Per Sec
MySQL Free Pages
MySQL Lazy Writes Per Sec
MySQL Lock Requests Per Se
MySQL Lock Timeouts Per S
MySQL Lock Wait Times
MySQL Lock Waits Per Sec
MySQL Log Cache Hit Rate

Now we have Network status map which is graphical representation of the network status



CONCLUSION : In this assignment we learned about how to continuously monitor Nagios commands and Installation and configuration of Nagios on Linux Machine.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 9

AIM : To understand Lambda Function and create a Lambda function which will log "An Image has been added" once you add an object to a specific bucket in S3.

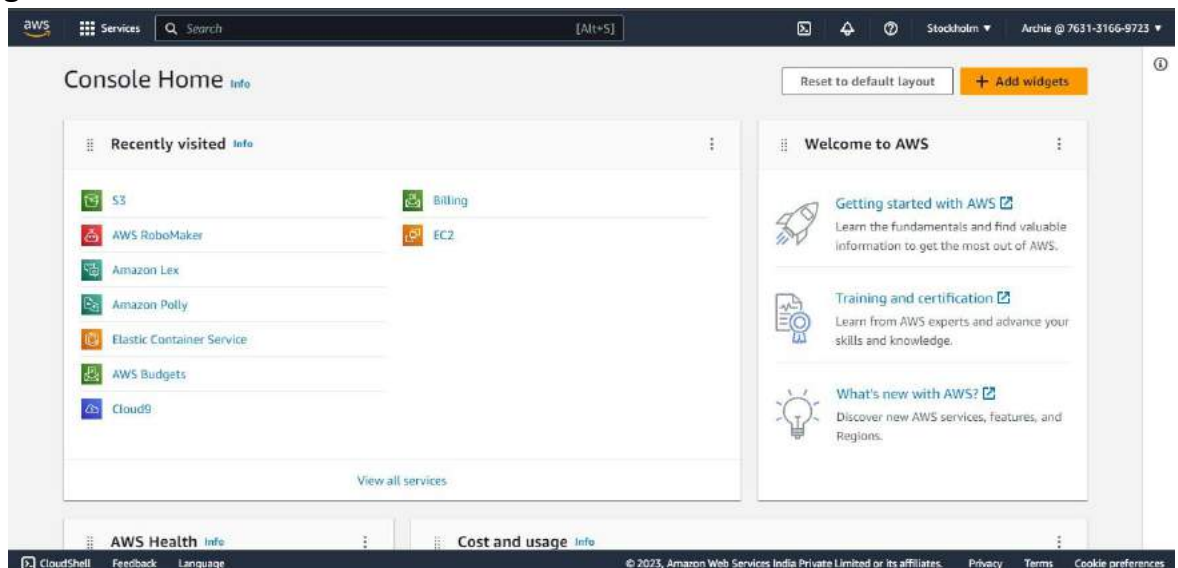
LO MAPPED : LO1,LO5

THEORY :

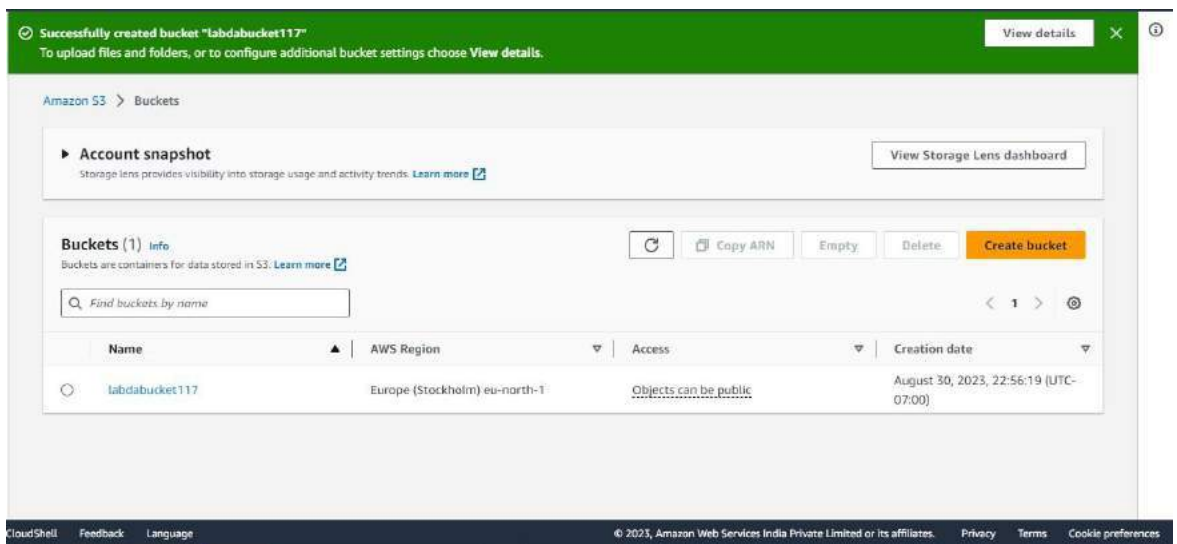
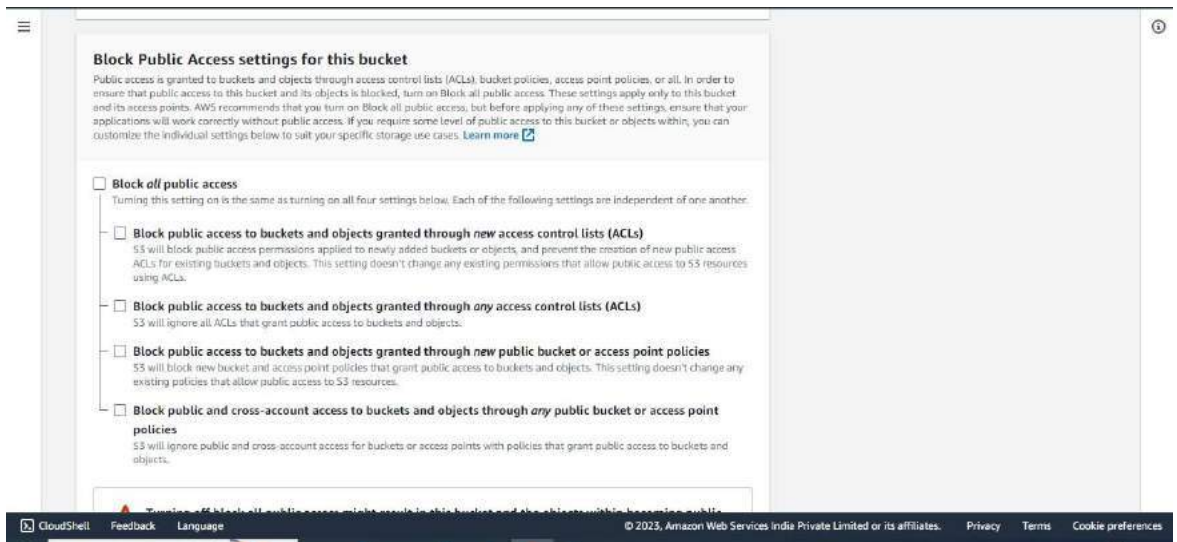
You can use Lambda to process event notifications from Amazon Simple Storage Service. Amazon S3 can send an event to a Lambda function when an object is created or deleted. You configure notification settings on a bucket, and grant Amazon S3 permission to invoke a function on the function's resource-based permissions policy.

STEPS TO FOLLOW:

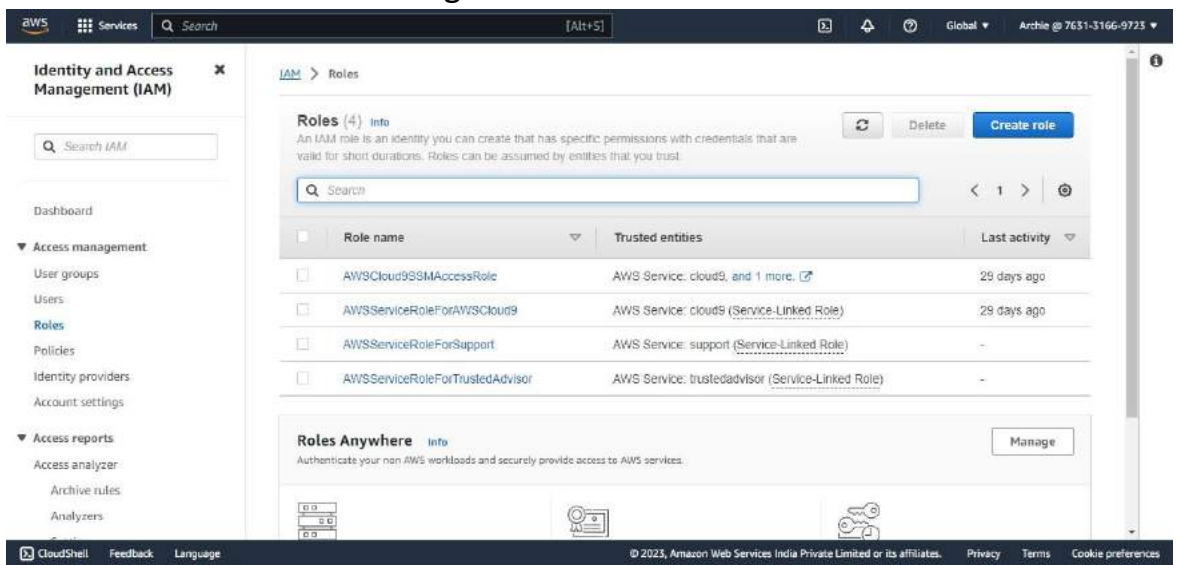
1. Log in as IAM User



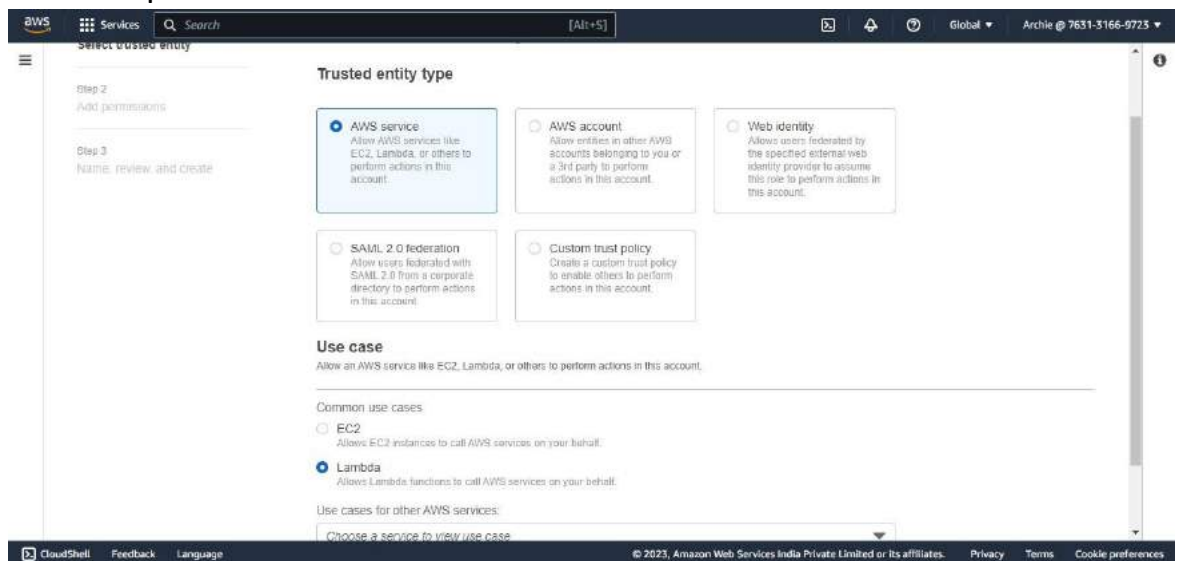
2. Create a S3 bucket and Enable the "Block all public access"



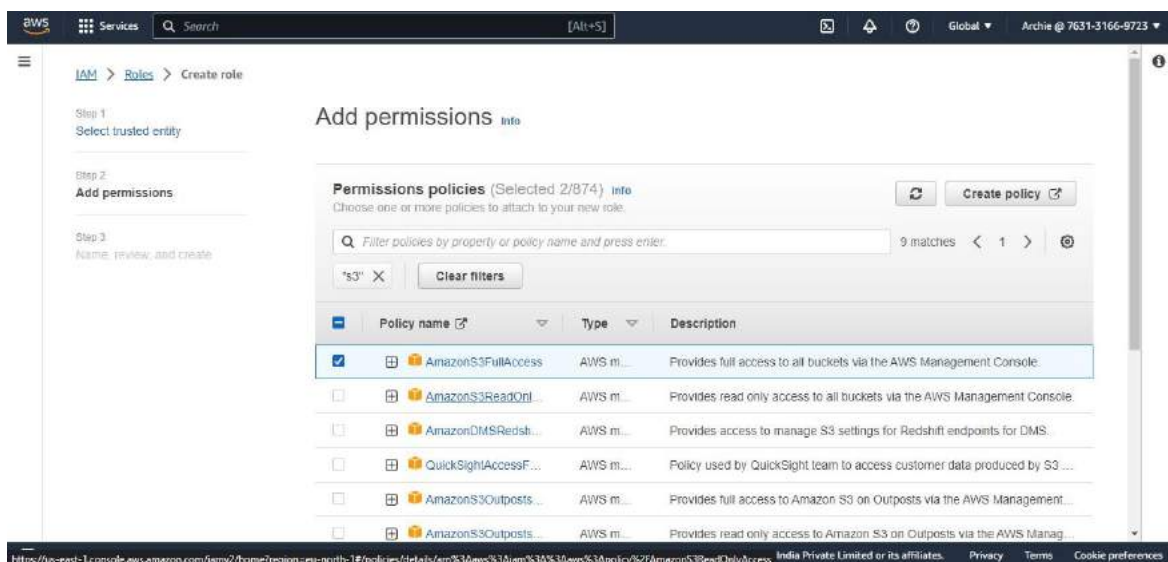
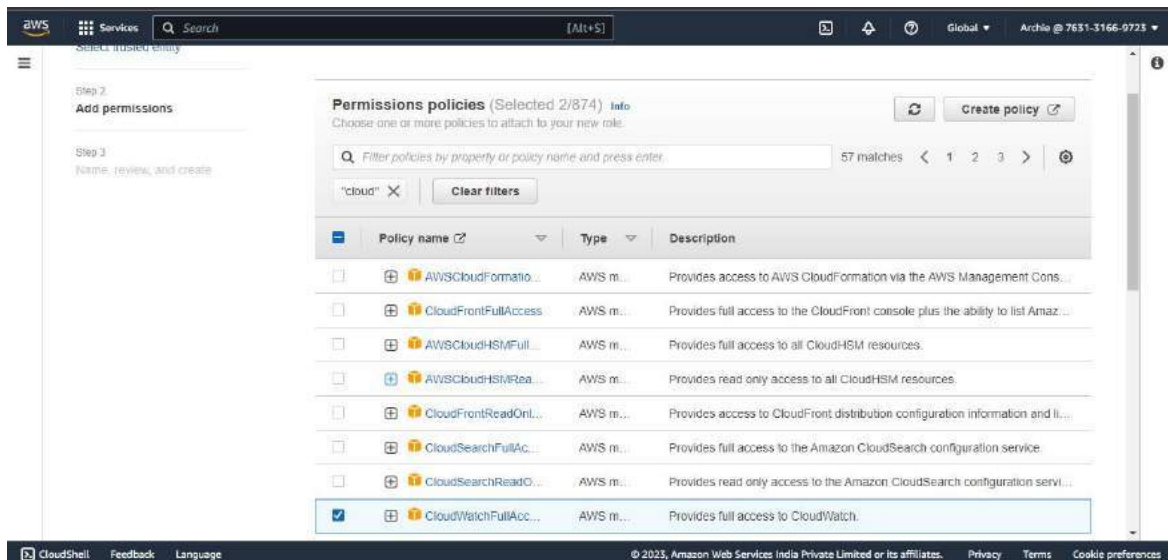
3. Search IAM on the console and go to "Roles". Click on Create Roles



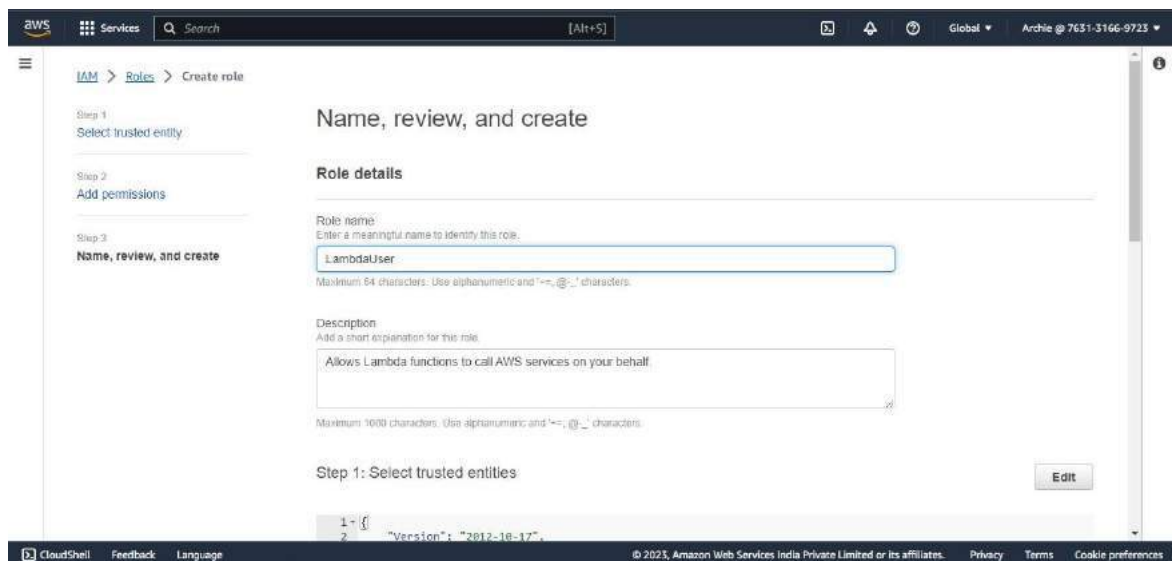
4. Select the options of "AWS service" and "lambda"



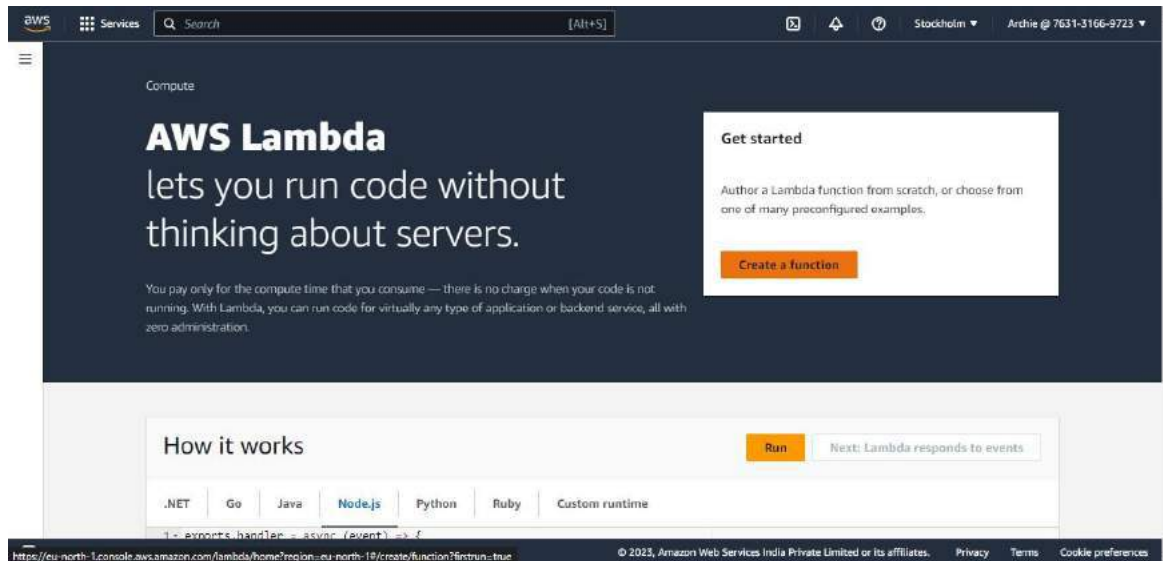
5. Enable the "CloudWatchFullAccess" and "AmazonS3FullAccess"



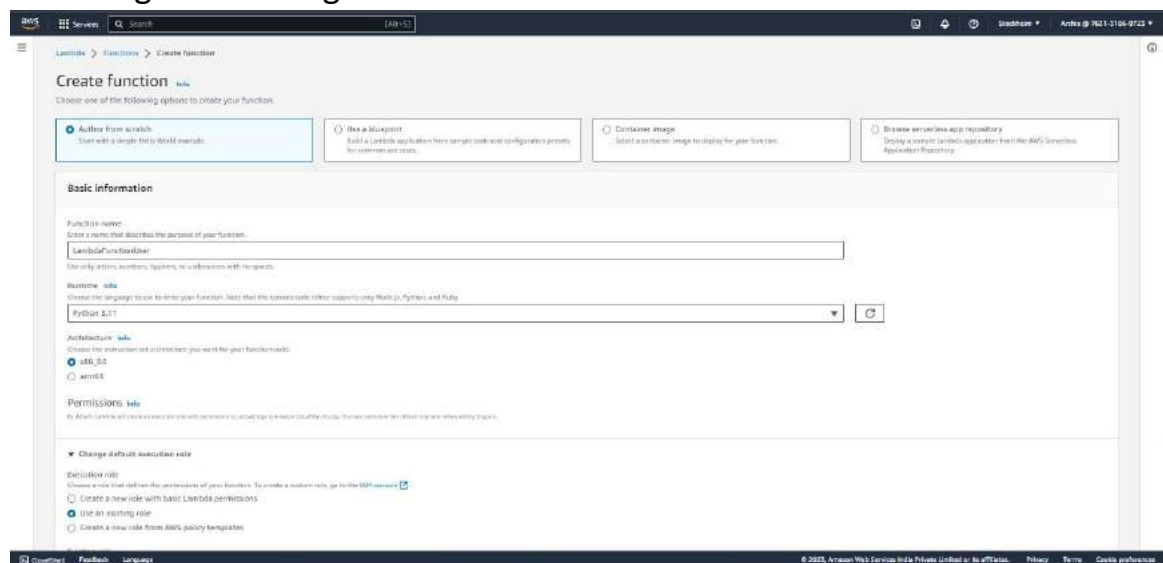
6. Click on Create Role and you will be redirected to this dashboard. Give the Role a Name and Click on Done.



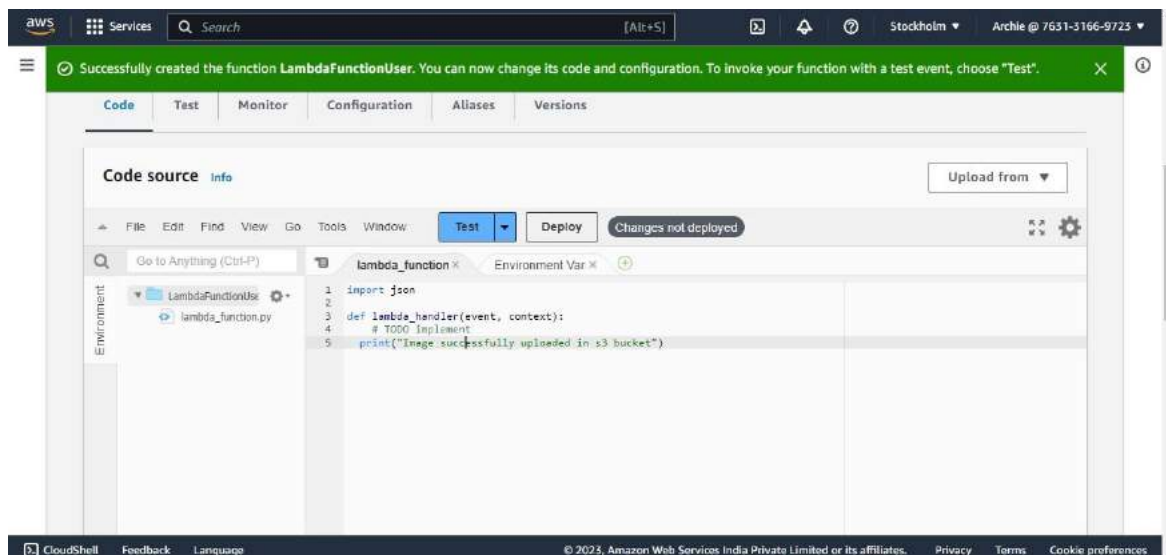
7. Search for Lambda in the console and click on Create Function.



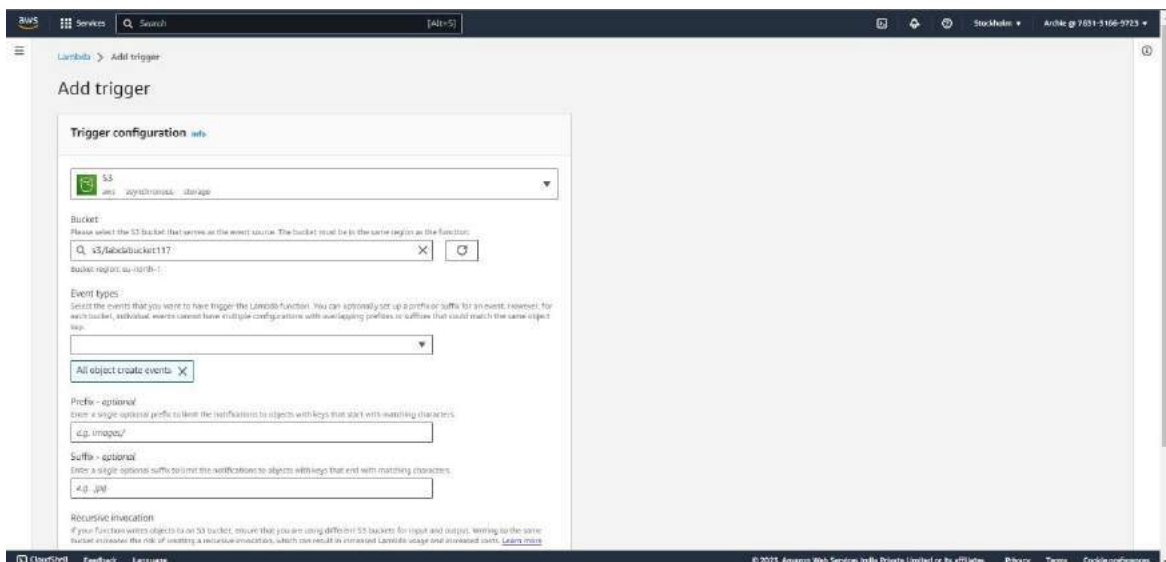
8. Change the settings of the Lambda Function.



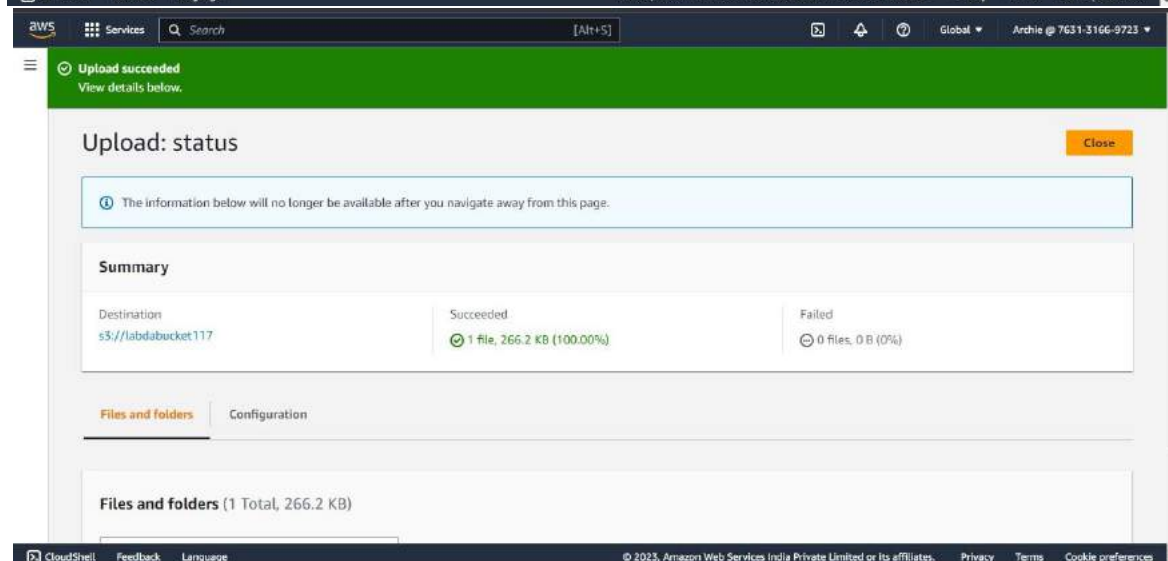
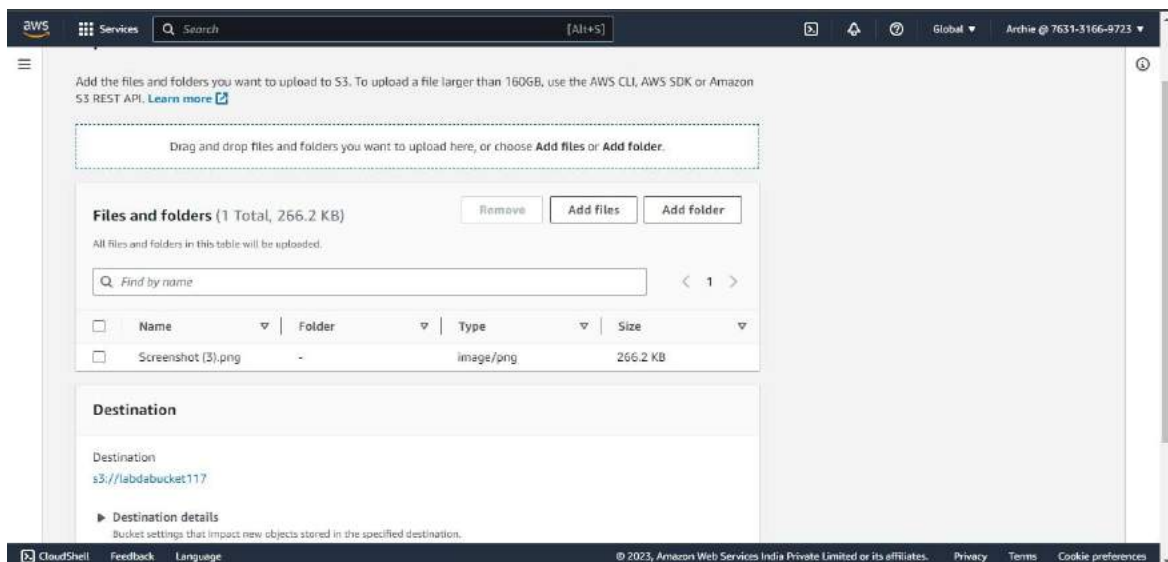
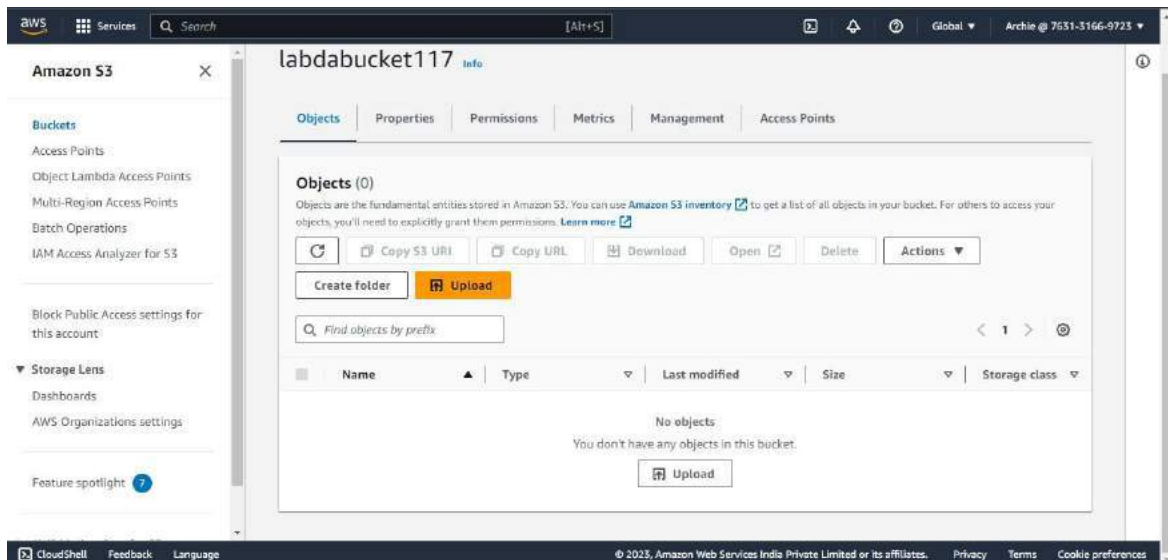
9. Add the python code and click on deploy to save the changes.



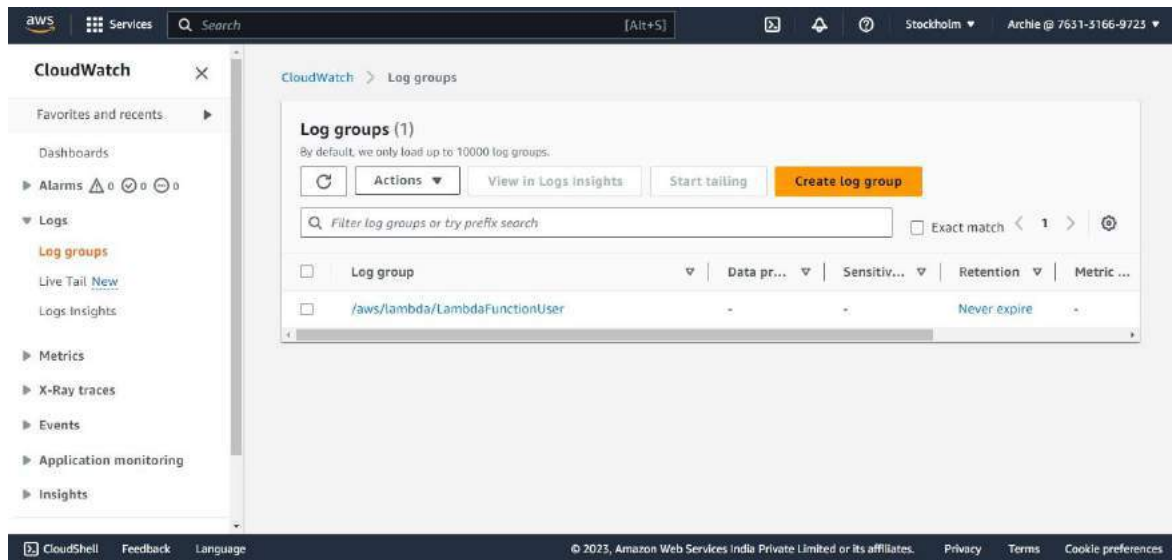
10. Scroll above and click on ADD TRIGGER. Select the following options and click on Done.



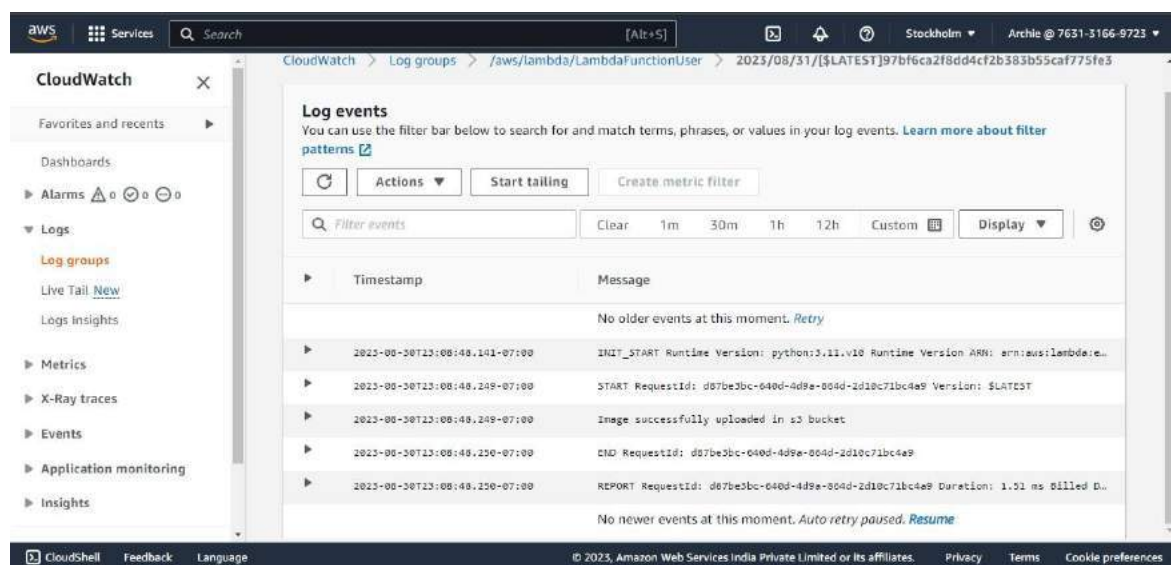
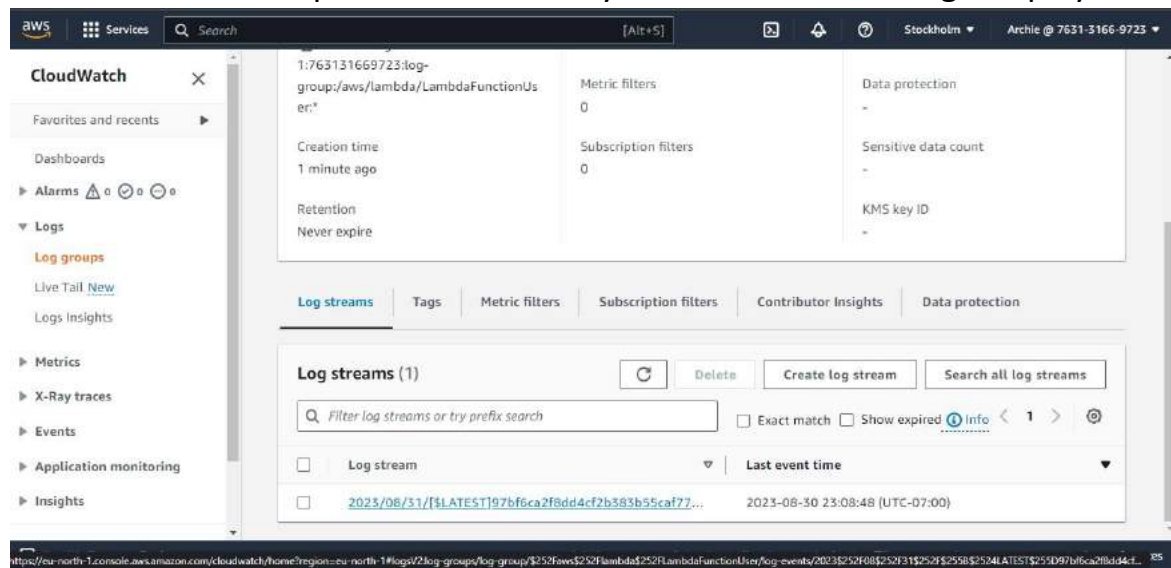
11. Go to S3 bucket and click on Add files. Select a image and click on Upload.



12. Search CloudWatch and go to Log groups. Select the existing Log Group.



13. Click on the link provided and then you will see the message displayed.



CONCLUSION : In this assignment, we learnt how to create a Lambda function which will log "An Image has been added" once you add an object to a specific bucket in S3.

NAME : SOUMIL SALVI

ROLL NO : 104

ASSIGNMENT 10

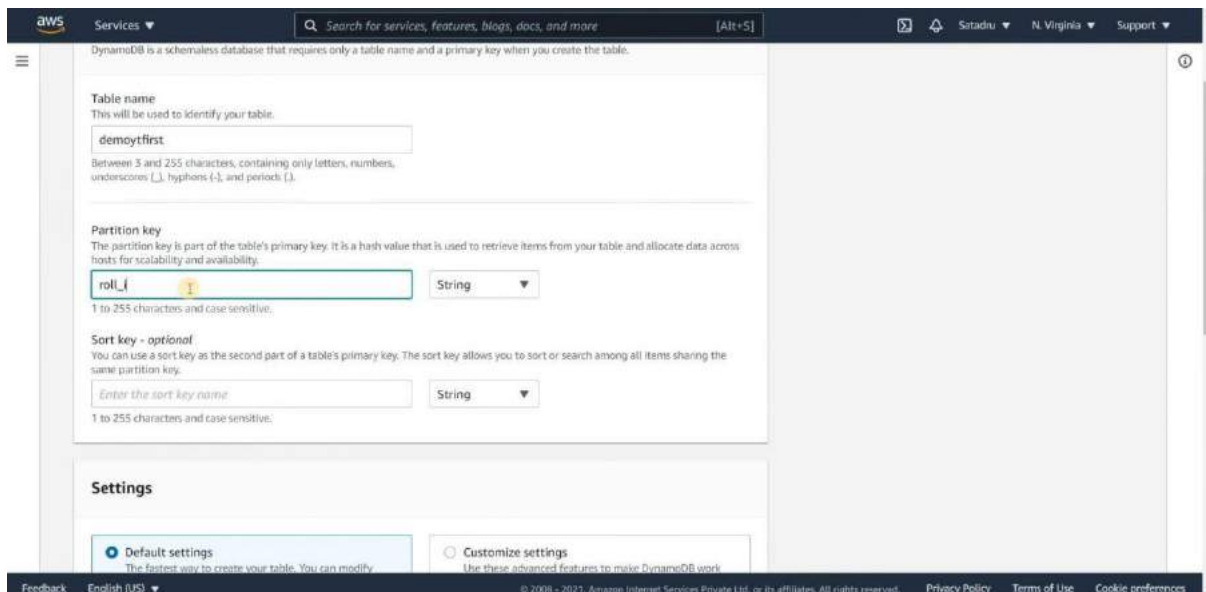
AIM : To create a Lambda function using Python for adding data to Dynamo DB database.

LO MAPPED : LO1,LO6

THEORY :

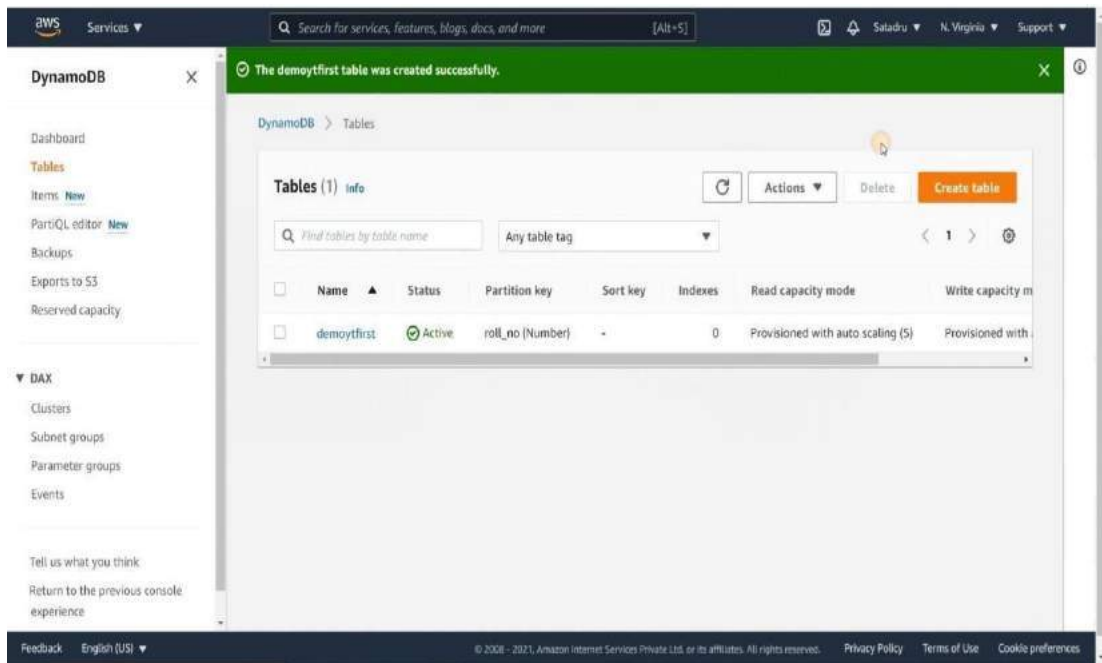
Step 1: Set Up AWS Environment

1. Create an AWS Account: If you don't have an AWS account, sign up for one at AWS Console.
2. Access AWS Lambda Console:



The screenshot shows the AWS Lambda console interface for creating a new function. The 'Table name' field is filled with 'demoytfirst'. The 'Partition key' field is filled with 'roll_id' and the data type is set to 'String'. The 'Sort key - optional' field is empty, and the data type is also set to 'String'. The 'Settings' section shows 'Default settings' selected.

- Go to the AWS Lambda Console.
 - Click on "Create function."
3. Create a New Lambda Function



- Choose "Author from scratch."
- Enter a name for your function (e.g., AddToDynamoDBFunction).
- Choose "Python" as the runtime.

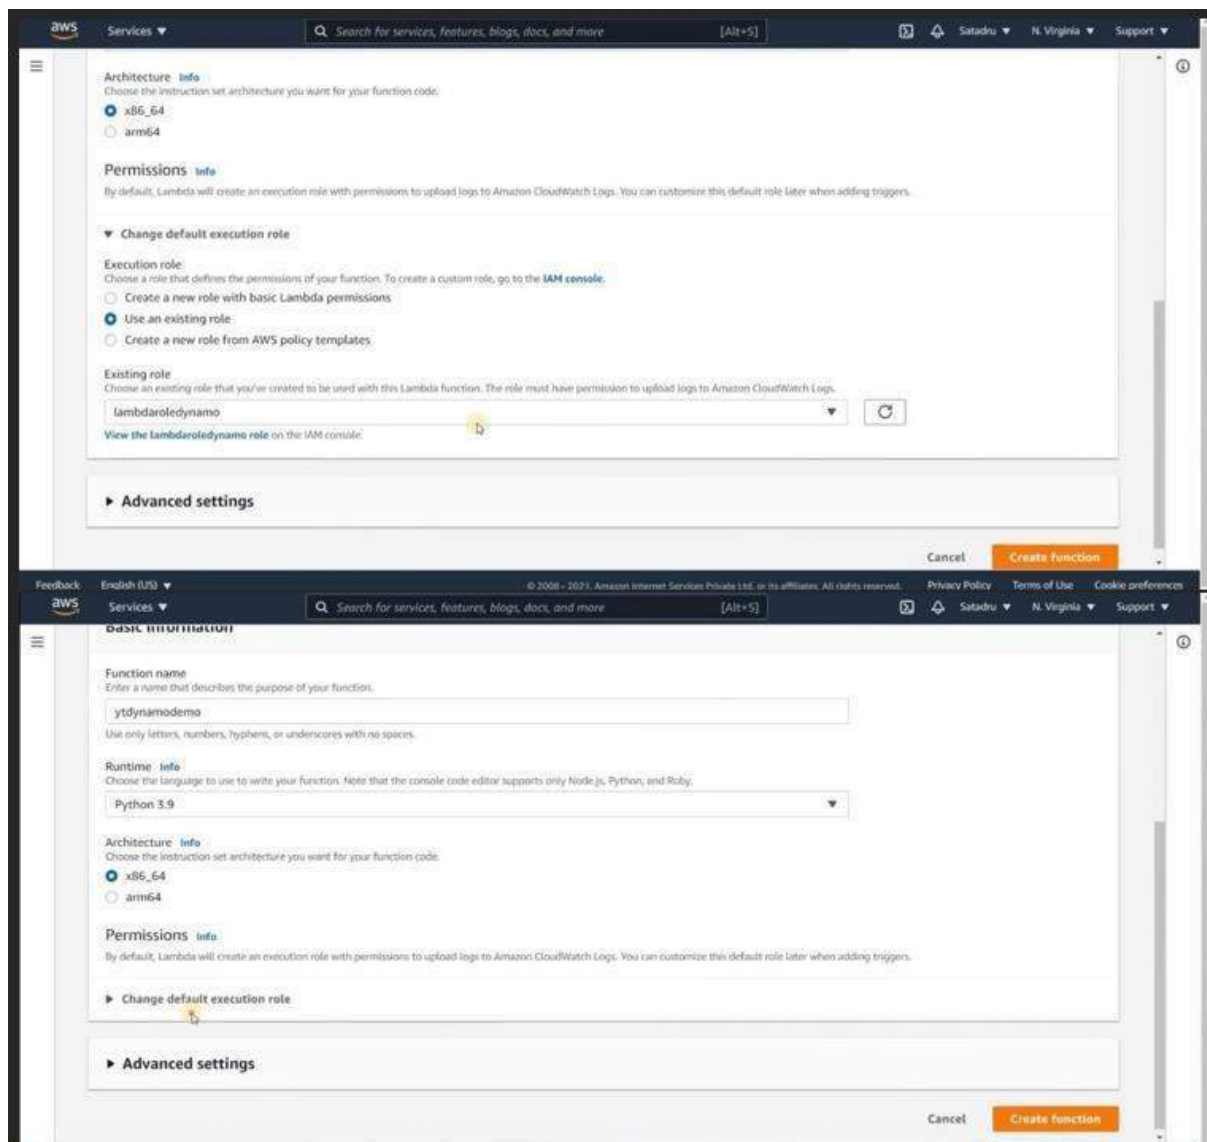
4. Configure Execution Role

- Create a new role with basic Lambda permissions and additional permissions to interact with DynamoDB.
- Attach the role to your Lambda function.

Step 2: Set Up DynamoDB

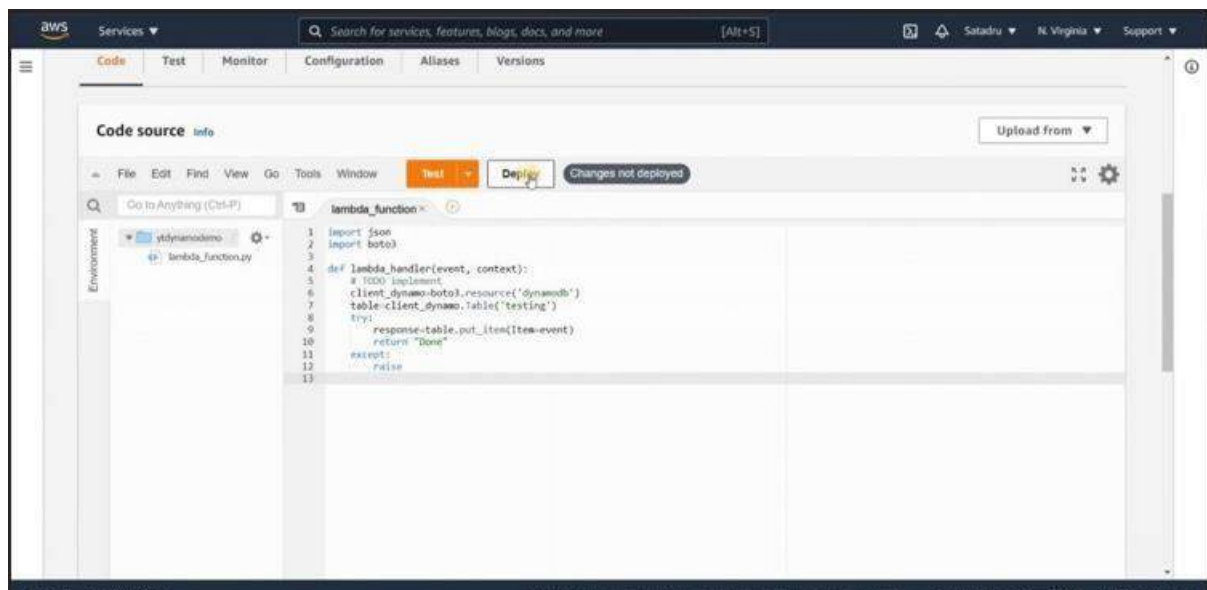
1. Access DynamoDB Console:

- Go to the AWS DynamoDB Console.

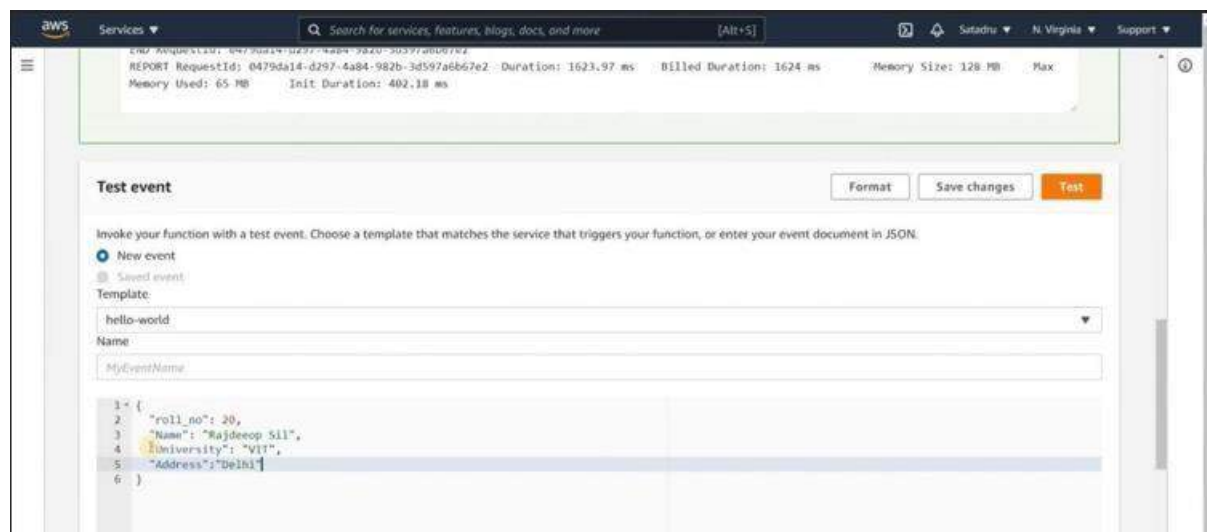


2. Create a DynamoDB Table:

- Click on "Create table."
- Enter a table name and a primary key (e.g., ID).
- Configure other settings as needed and create the table.



Step 3: Write Lambda Function Code



. Configure Lambda Handler:

- In the Lambda function configuration, set the handler to filename.lambda_handler (e.g., yourfilename.lambda_handler).

Step 4: Configure Environment Variables

1. Add DynamoDB Table Name:

- In the Lambda function configuration, add an environment variable (e.g., DYNAMODB_TABLE) with the value set to your DynamoDB table name.

Step 5: Test Locally

1. Test Lambda Function Locally:

- Create a test event with sample data.
- Test the Lambda function using the "Test" button in the Lambda console.

Step 6: Deploy Lambda Function

1. Deploy the Lambda Function:

- Click on the "Deploy" button in the Lambda console.

Step 7: Test in AWS Lambda Console

1. Test in AWS Lambda Console:

- Use the Lambda console to test the function by creating a test event with sample data.
- Verify that the function executes successfully.

Step 8: Monitor and Troubleshoot

1. Check CloudWatch Logs:

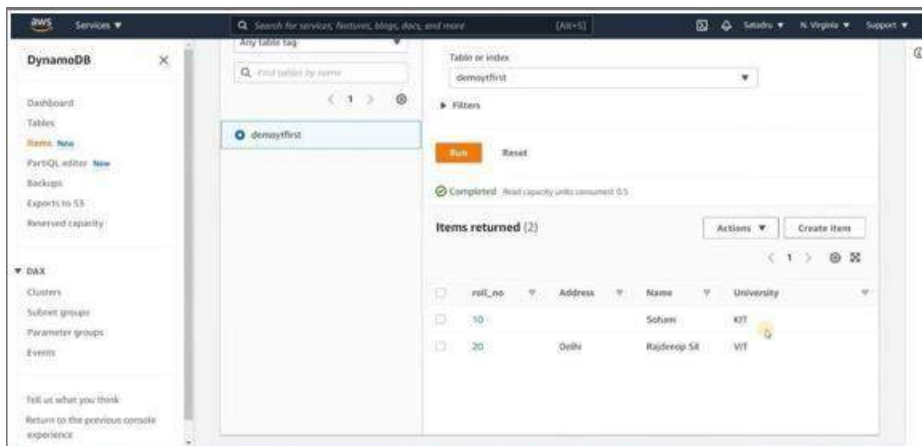
- Use CloudWatch Logs to check logs generated by your Lambda function for troubleshooting.

Step 9: Documentation

1. Create Documentation:

- Document the purpose, input parameters, and expected output of your Lambda function.
- Explain any challenges faced during development and how they were resolved.

By following these steps, you can create a Lambda function using Python to add data to a DynamoDB database in the AWS environment.



CONCLUSION : By this assignment we learn how to create a Lambda function using Python for adding data to Dynamo DB database.

NAME : SOUMIL SALVI

ROLL NO : 104

WRITTEN ASSIGNMENT - 1

1. what security measures can be taken while using Kubernetes?

- i. **Role-Based Access Control (RBAC):** RBAC restricts who can perform actions within a Kubernetes cluster. It defines roles and role bindings to specify what resources and operations users or service accounts can access. This prevents unauthorized access and actions within the cluster.
- ii. **Regular Updates:** Keeping Kubernetes and its components up to date is crucial. New releases often include security patches. Regular updates help mitigate known vulnerabilities and ensure your cluster remains secure.
- iii. **Network Policies:** Network policies allow you to define rules for communication between pods. By specifying which pods can communicate with each other, you can limit the attack surface and prevent unauthorized access.
- iv. **Container Security Tools:** Employ container security tools like vulnerability scanners to assess the security of container images. These tools can identify and remediate vulnerabilities in the containerized applications before they are deployed.
- v. **Monitoring and Audit:** Implement monitoring and auditing solutions to track cluster activity. This helps detect and respond to suspicious or unauthorized behavior. Tools like Prometheus and Grafana can be used for monitoring, while audit logs provide insights into cluster activity.
- vi. **Secrets Management:** Sensitive data like API keys, passwords, and certificates should be stored securely using Kubernetes secrets or external vaults. This prevents sensitive information from being exposed within

containers or configuration files.

- vii. PodSecurityPolicies (PSP): PSP is a Kubernetes feature that enforces security policies at the pod level. It allows you to define restrictions on privilege escalation, host access, and other security-sensitive configurations for pods.
- viii. Namespaces: Use Kubernetes namespaces to logically isolate workloads. This provides a level of separation between different applications or teams, reducing the risk of unauthorized access or interference between them.
- ix. Admission Controllers: Admission controllers are webhook plugins that intercept and validate requests to the Kubernetes API server. You can use them to enforce custom policies and ensure that only compliant resources are admitted to the cluster.
- x. Container Runtime Security: Implement container runtime security solutions like Docker Security Scanning or container runtime protection tools. These tools monitor containers at runtime for abnormal behavior, helping to detect and respond to potential threats.

2. What are the three security techniques that can be used to protect data?

Three security techniques commonly used to protect data are:

- i. Encryption: Encryption is the process of converting data into a secure format that can only be read by someone with the decryption key. It ensures that even if unauthorized parties access the data, they cannot understand it without the correct key. Two common types of encryption are:
 - a. Data-at-rest Encryption: Protects data when it's stored on disk or in a database.
 - b. Data-in-transit Encryption: Secures data as it's transmitted between systems over networks.

- ii. Access Control: Access control mechanisms regulate who can access data and what actions they can perform on it. This involves setting permissions, roles, and policies to ensure that only authorized users or applications can access and manipulate data. Role-Based Access Control (RBAC) and Attribute-Based Access Control (ABAC) are commonly used access control models.
- iii. Data Masking/Redaction: Data masking or redaction involves obscuring or replacing sensitive data with fictitious or scrambled values. This is often used in non-production environments or when sharing data with third parties. It ensures that even if someone gains access to the data, they cannot see the actual sensitive information.

These techniques are often used in combination to create a layered approach to data security, providing multiple levels of protection to safeguard sensitive information from unauthorized access and disclosure.

3. How do you expose a service using ingress in Kubernetes?

To expose a service using Ingress in Kubernetes, you need to follow these steps:

- i. **Set up Kubernetes**: Ensure you have a Kubernetes cluster up and running, and you have the `kubectl` command-line tool configured to communicate with the cluster.
- ii. **Deploy Your Application**: Deploy your application as a Kubernetes Deployment or a Pod, and create a Kubernetes Service to expose it internally within the cluster. This Service will be the target for the Ingress.
- iii. **Install an Ingress Controller**: You need to have an Ingress controller installed in your cluster. Some popular options include Nginx Ingress Controller, Traefik, or HAProxy Ingress. The controller will manage the Ingress resources and configure the load balancer.

For example, to install the Nginx Ingress Controller, you can use:

```
```bash
```

```
kubectl apply -f
https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.0.0/deploy/static/provider/cloud/deploy.yaml
```

```
...
```

- iv. Create an Ingress Resource: Define an Ingress resource that specifies the rules for routing traffic to your service. Here's an example Ingress resource manifest:

```
```yaml
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: my-ingress
```

```
spec:
```

```
  rules:
```

```
    - host: example.com
```

```
      http:
```

```
        paths:
```

```
          - path: /path
```

```
            pathType: Prefix
```

```
            backend:
```

```
              service:
```

```
                name: your-service
```

```
                port:
```

```
                  number: 80
```

```
```
```

In this example, traffic for `example.com/path` will be routed to `your-service`.

- v. Apply the Ingress Resource: Use `kubectl apply` to create the Ingress resource in your cluster:

```
```bash
```

```
kubectl apply -f your-ingress.yaml
```

```
```
```

- vi. Configure DNS: Ensure that the DNS records for the specified hostname (e.g., `example.com`) point to the external IP address of your Ingress controller.
- vii. Access Your Service: After DNS propagation, you should be able to access your service externally via the hostname and path you defined in the Ingress resource.

#### 4. Which service protocols does Kubernetes ingress expose?

Kubernetes Ingress is primarily designed to expose HTTP and HTTPS services, making it suitable for routing and load balancing web traffic. However, with the evolution of Kubernetes and Ingress controllers, it has expanded to support additional protocols and features:

- i. **HTTP:** Ingress is commonly used to expose HTTP services. You can define routing rules based on URL paths, hostnames, and other HTTP attributes.
- ii. **HTTPS:** Secure HTTP services can be exposed through Ingress by configuring TLS certificates. This allows you to terminate SSL/TLS encryption at the Ingress controller and route decrypted traffic to your services.
- iii. **TCP:** Some Ingress controllers, like Nginx Ingress, support TCP services. This enables you to expose non-HTTP services such as databases or custom protocols. TCP-based routing typically relies on port numbers.
- iv. **UDP:** While less common, some Ingress controllers support UDP services.

UDP is a connectionless protocol used for various purposes, including DNS and VoIP. Exposing UDP services may require specific controller support.

- v. **gRPC:** If your services use the gRPC protocol, you can configure Ingress resources to handle gRPC traffic. gRPC is a high-performance RPC (Remote Procedure Call) framework often used for communication between microservices.
- vi. **WebSocket:** Ingress controllers can be configured to support WebSocket connections. WebSocket is a protocol that enables full-duplex communication over a single TCP connection and is used for real-time applications.
- vii. **Custom Protocols:** In some cases, you may need to expose services using custom or proprietary protocols. Depending on your Ingress controller and its capabilities, you might be able to configure it to handle these custom protocols.

Additionally, Ingress controllers often evolve, so it's essential to refer to the documentation and features of the specific controller you plan to use to ensure compatibility with your service protocols.

NAME : SOUMIL SALVI

ROLL NO : 104

## WRITTEN ASSIGNMENT 2

### Q1. How to deploy lambda function on AWS?

Deploying a Lambda function on AWS involves several steps. Lambda is a serverless compute service that lets you run code without provisioning or managing servers. Here's a step-by-step guide on how to deploy a Lambda function:

#### Step 1: Create a Lambda Function

1. Log in to the AWS Management Console.
2. Navigate to the Lambda service by searching for it in the AWS Services section.
3. Click the "Create function" button.

#### Step 2: Configure Your Function

4. Choose "Author from scratch."
5. Fill in the basic information for your function, including the name, runtime, and execution role.
6. For "Execution role," you can create a new role from a template or use an existing role if you have one with the necessary permissions.
7. Click the "Create function" button.

#### Step 3: Upload Your Code

8. In the "Function code" section, you can upload your code either directly (ZIP file or folder) or by specifying a repository from AWS CodeCommit, Amazon S3, or other options.
9. Configure the handler if needed. The handler is the method that AWS Lambda invokes.



10. Set environment variables if your code requires them.
11. Adjust the runtime settings if necessary.
12. Click the "Deploy" button to upload your code.

#### **Step 4: Define the Trigger**

13. In the "Add triggers" section, you can specify event sources that will trigger your Lambda function. This can be an API Gateway, an S3 bucket, an SNS topic, etc.
14. Configure the trigger according to your needs and grant permissions as required.

#### **Step 5: Test Your Function**

15. You can test your Lambda function by creating a test event or using a sample test event provided by AWS.
16. Click the "Test" button to run your function and see the results.

#### **Step 6: Monitor and Troubleshoot (Optional)**

17. AWS Lambda provides various monitoring and logging options. You can configure CloudWatch alarms, inspect logs, and set up notifications to keep an eye on your function's performance.

#### **Step 7: Save and Deploy the Function**

18. After you've configured everything to your satisfaction, click the "Save" button to save your changes.
19. Click the "Deploy" button to make your function live and ready to respond to triggers.

#### **Step 8: Invocation and Scaling**

Your Lambda function is now deployed and ready to be invoked by the trigger you configured. AWS Lambda handles the scaling of resources based on the incoming workload automatically.

## **Q2. What are the deployment options for AWS Lambda?**

AWS Lambda offers several deployment options:

- 1. Code Upload:** You can directly upload your function code as a ZIP file or a deployment package when creating or updating your Lambda function.
- 2. AWS Lambda Layers:** You can use Lambda Layers to separate your function code from its dependencies. This allows you to manage and version common libraries independently and share them across multiple functions.
- 3. AWS SAM (Serverless Application Model):** AWS SAM is a framework for building serverless applications. You can define your Lambda functions and their associated resources in a SAM template file, then deploy the entire application using AWS SAM CLI.
- 4. AWS CloudFormation:** You can use AWS CloudFormation templates to define and deploy Lambda functions along with other AWS resources. This enables infrastructure as code (IaC) for your serverless applications.
- 5. AWS Serverless Application Repository:** You can publish and share serverless applications and Lambda functions using the AWS Serverless Application Repository. Users can deploy your applications directly from the repository.
- 6. AWS CodePipeline:** You can integrate AWS Lambda deployment into a continuous integration/continuous deployment (CI/CD) pipeline using AWS CodePipeline. This automates the building, testing, and deployment of your Lambda functions.
- 7. Container Image Support:** AWS Lambda now supports deploying functions as container images, allowing you to package your function and dependencies in a Docker container and deploy them as a Lambda function.

These deployment options provide flexibility and enable you to choose the one that best fits your application architecture and development workflow.

### **Q3 What are the 3 full deployment modes that can be used for AWS?**

In the context of AWS Lambda, there are three primary deployment modes:

**1. Automatic Deployment:** AWS Lambda provides automatic deployment when you directly upload your function code as a ZIP file or specify a deployment package. This is the most common and straightforward deployment method, and it's suitable for most use cases.

**2. Infrastructure as Code (IaC) with CloudFormation:** AWS CloudFormation is a service that allows you to define and provision AWS infrastructure and resources in a template. You can use CloudFormation to define your Lambda functions, their associated resources, and any other AWS resources your application needs. When you create or update the CloudFormation stack, it deploys the Lambda functions and other resources defined in the template. This approach is more suitable for complex serverless applications and infrastructure orchestration.

**3. Serverless Application Model (AWS SAM):** AWS SAM is an open-source framework for building serverless applications. It extends AWS CloudFormation to provide a simplified way to define serverless resources, including Lambda functions. You can use AWS SAM to define your functions and related resources in a SAM template, and then deploy the entire application using AWS SAM CLI. This approach is a balance between the simplicity of direct deployment and the power of CloudFormation for more complex applications.

These deployment modes offer different levels of control and abstraction, allowing you to choose the one that best matches your deployment needs and development workflow.

### **Q4. What are the 3 components of AWS Lambda?**

AWS Lambda consists of three primary components:

**1. Function Code:** This is the core component of AWS Lambda. It's the code that you want to run in response to events. You can write your code in various supported programming languages (e.g., Node.js, Python, Java, C#, etc.). You can package your code along with its dependencies into a deployment package, which you upload to Lambda.

**2. Event Sources:** Event sources are triggers that invoke your Lambda function. AWS Lambda supports various event sources, such as Amazon S3, Amazon DynamoDB, Amazon SNS, AWS API Gateway, and more. When an event occurs in one of these services, it triggers the execution of your Lambda function.

**3. Execution Environment:** AWS Lambda manages the execution environment for your function. It automatically scales and provisions the necessary infrastructure to run your code. You don't need to worry about server provisioning or resource management. AWS Lambda also provides runtime support for various programming languages, so your code runs in an isolated environment with the necessary resources to execute.